

Fatima Jinnah Women University



Course Title : Cloud Computing
Assignment 2

Submitted by :

Lujain Zia (2023-BSE-034)

Submitted to :

Sir Waqas Saleem

Date : December 28, 2025

Part 1:
Infrastructure Setup (25 marks)

1.1 Project Structure (5 marks)

Create a well-organized Terraform project with the following structure:

Assignment2/

```
└── main.tf
└── variables.tf
└── outputs.tf
└── locals.tf
└── terraform.tfvars
└── .gitignore
└── modules/
    ├── networking/
    │   ├── main.tf
    │   ├── variables.tf
    │   └── outputs.tf
    ├── security/
    │   ├── main.tf
    │   ├── variables.tf
    │   └── outputs.tf
    └── webserver/
        ├── main.tf
        ├── variables.tf
        └── outputs.tf
└── scripts/
    ├── nginx-setup.sh
    └── apache-setup.sh
└── README.md
```

Tasks:

- Create all necessary files and directories
- Implement proper .gitignore to exclude sensitive files
- Document the project structure in README.md

Deliverables:

Screenshot: assignment_part1_project_structure.png (tree command output)

```
C:.
  .gitignore
  locals.tf
  main.tf
  outputs.tf
  README.md
  terraform.tfvars
  variables.tf

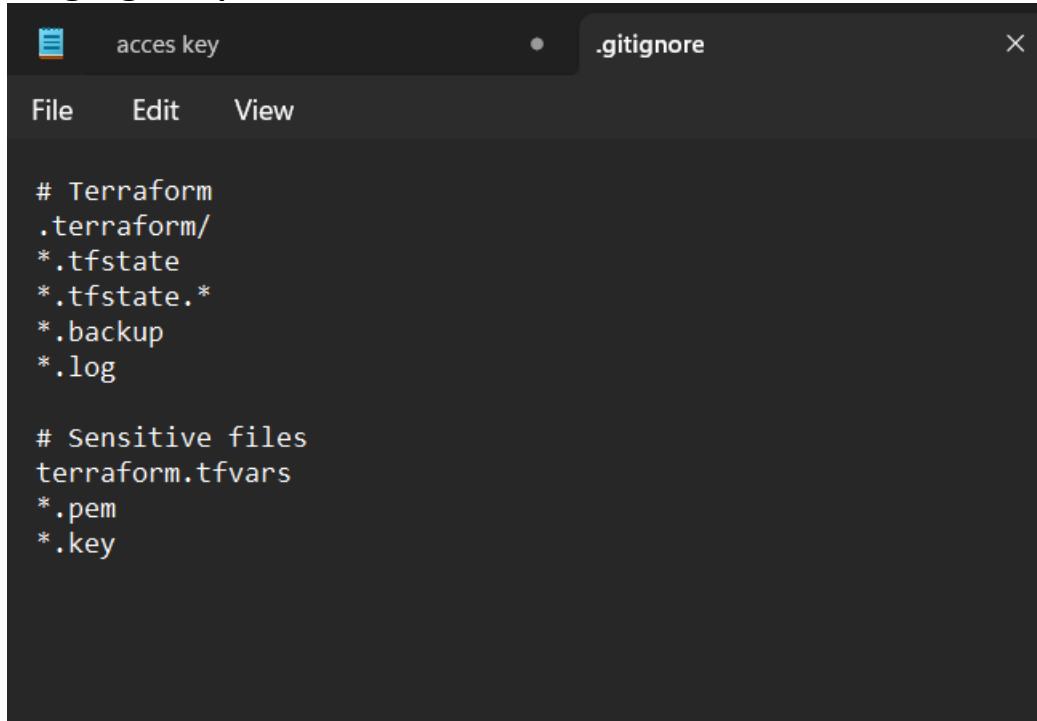
  modules
    networking
      main.tf
      outputs.tf
      variables.tf

    security
      main.tf
      outputs.tf
      variables.tf

    webserver
      main.tf
      outputs.tf
      variables.tf

  scripts
    apache-setup.sh
    nginx-setup.sh
```

Screenshot: assignment_part1_gitignore.png (content of .gitignore)



The screenshot shows a code editor window with the title bar "acces key" and the file name ".gitignore". The menu bar includes "File", "Edit", and "View". The content of the file is as follows:

```
# Terraform
.terraform/
*.tfstate
*.tfstate.*
*.backup
*.log

# Sensitive files
terraform.tfvars
*.pem
*.key
```

1.2 Variable Configuration (5 marks)

Define all required variables in variables.tf:

Required Variables:

vpc_cidr_block (string)
subnet_cidr_block (string)

availability_zone (string)
env_prefix (string)
instance_type (string)
public_key (string)
private_key (string)
backend_servers (list of objects with name and script_path)

Tasks:

Add validation rules for CIDR blocks

Add descriptions for all variables

Set appropriate defaults where applicable

Populate terraform.tfvars with your values

Sample terraform.tfvars structure:

```
vpc_cidr_block  = "10.0.0.0/16"subnet_cidr_block =
"10.0.10.0/24"availability_zone = "me-central-1a"env_prefix      =
"prod"instance_type   = "t3.micro"public_key      =
"~/.ssh/id_ed25519.pub"private_key     = " ~/.ssh/id_ed25519"
```

Deliverables:

Screenshot: assignment_part1_variables_tf.png

```
variable "vpc_cidr_block" {
  description = "CIDR block for the VPC"
  type        = string
  validation {
    condition  = can(regex("^(0-9){1,3}\.(0-9){1,3}\.(0-9){1,3}\.(0-9){1,3}/[0-9]+$", var.vpc_cidr_block))
    error_message = "Invalid CIDR format for VPC"
  }
}

variable "subnet_cidr_block" {
  description = "CIDR block for the subnet"
  type        = string
  validation {
    condition  = can(regex("^(0-9){1,3}\.(0-9){1,3}\.(0-9){1,3}\.(0-9){1,3}/[0-9]+$", var.subnet_cidr_block))
    error_message = "Invalid CIDR format for Subnet"
  }
}

variable "availability_zone" {
  description = "AWS availability zone (e.g. me-central-1a)"
  type        = string
}

variable "env_prefix" {
  description = "Environment name (e.g. dev, prod)"
  type        = string
  default     = "dev"
}

variable "instance_type" {
```

Screenshot: assignment_part1_terraform_tfvars.png

```

File Edit View

vpc_cidr_block    = "10.0.0.0/16"
subnet_cidr_block = "10.0.10.0/24"
availability_zone = "me-central-1a"
env_prefix        = "prod"
instance_type     = "t3.micro"
public_key        = "C:/Users/user/.ssh/id_ed25519.pub"
private_key       = "C:/Users/user/.ssh/id_ed25519"

backend_servers = [
  {
    name      = "web-1"
    script_path = "./scripts/apache-setup.sh"
  },
  {
    name      = "web-2"
    script_path = "./scripts/apache-setup.sh"
  },
  {
    name      = "web-3"
    script_path = "./scripts/apache-setup.sh"
  }
]

```

1.3 Networking Module (5 marks)

Create a networking module that provisions:

VPC with specified CIDR block

Subnet with public IP assignment enabled

Internet Gateway

Route table with default route to IGW

Associate route table with subnet

Module Location: modules/networking/

Required Outputs:

vpc_id
subnet_id
igw_id
route_table_id

Tasks:

- Create VPC resource
- Create subnet with map_public_ip_on_launch = true
- Create and attach Internet Gateway
- Configure routing table
- Add proper tags to all resources using env_prefix

Deliverables:

Screenshot: assignment_part1_networking_module_main.png

```

resource "aws_vpc" "main" {
  cidr_block = var.vpc_cidr_block
  tags = {
    Name = "${var.env_prefix}-vpc"
  }
}

resource "aws_subnet" "main" {
  vpc_id      = aws_vpc.main.id
  cidr_block   = var.subnet_cidr_block
  availability_zone = var.availability_zone
  map_public_ip_on_launch = true

  tags = {
    Name = "${var.env_prefix}-subnet"
  }
}

resource "aws_internet_gateway" "main" {
  vpc_id = aws_vpc.main.id

  tags = {
    Name = "${var.env_prefix}-igw"
  }
}

resource "aws_route_table" "main" {
  vpc_id = aws_vpc.main.id

  route {
    ...
  }
}

```

Screenshot: assignment_part1_networking_module_outputs.png

g

```

aws_key_pair.ssh_key: Creating...
aws_vpc.myapp_vpc: Creating...
aws_key_pair.ssh_key: Creation complete after 1s [id=serverkey]
aws_vpc.myapp_vpc: Creation complete after 2s [id=vpc-0e27d646a98e5e382]
aws_internet_gateway.myapp_igw: Creating...
aws_subnet.myapp_subnet_1: Creating...
aws_default_security_group.myapp_sg: Creating...
aws_internet_gateway.myapp_igw: Creation complete after 0s [id=igw-04f131a91c7192065]
aws_route_table.myapp_route_table: Creating...
aws_subnet.myapp_subnet_1: Creation complete after 1s [id=subnet-0f568f6c021d0dfb9]
aws_route_table.myapp_route_table: Creation complete after 2s [id=rtb-00a3481e95c03304e]
aws_route_table_association.myapp_route_table_assoc: Creating...
aws_route_table_association.myapp_route_table_assoc: Creation complete after 0s [id=rtbassoc-088799ddd9d1d75cf]
aws_default_security_group.myapp_sg: Creation complete after 2s [id=sg-07c94d09956adbc84]
aws_instance.myapp-server: Creating...
aws_instance.myapp-server: Still creating... [00m10s elapsed]
aws_instance.myapp-server: Creation complete after 13s [id=i-0ad257af3bed851d5]

Apply complete! Resources: 8 added, 0 changed, 0 destroyed.

```

1.4 Security Module (5 marks)

Create a security module that provisions:

Two Security Groups:

Nginx Security Group (for reverse proxy/load balancer):

Ingress: Port 22 (SSH) from your IP only

Ingress: Port 80 (HTTP) from anywhere (0.0.0.0/0)

Ingress: Port 443 (HTTPS) from anywhere (0.0.0.0/0)

Egress: All traffic

Backend Security Group (for web servers):

Ingress: Port 22 (SSH) from your IP only

Ingress: Port 80 (HTTP) from Nginx security group only

Egress: All traffic

Module Location: modules/security/

Required Variables:

vpc_id

env_prefix

my_ip

Required Outputs:

nginx_sg_id

backend_sg_id

Tasks:

Create Nginx security group with appropriate rules

Create backend security group with appropriate rules

Use security group IDs for backend ingress (not CIDR blocks)

Add descriptive names and tags

Deliverables:

Screenshot: assignment_part1_security_module. Png

```
File Edit View

resource "aws_security_group" "nginx_sg" {
  name        = "${var.env_prefix}-nginx-sg"
  description = "Security group for NGINX load balancer"
  vpc_id      = var.vpc_id

  ingress {
    description = "Allow SSH from my IP"
    from_port   = 22
    to_port     = 22
    protocol   = "tcp"
    cidr_blocks = [var.my_ip]
  }

  ingress {
    description = "Allow HTTP from anywhere"
    from_port   = 80
    to_port     = 80
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    description = "Allow HTTPS from anywhere"
    from_port   = 443
    to_port     = 443
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

Screenshot: assignment_part1_security_groups_console.png (AWS Console)

Security Groups (5) Info					
<input type="text"/> Find security groups by attribute or tag					
	Name	Security group ID	Security group name	VPC ID	Description
<input type="checkbox"/>	prod-nginx-sg	sg-01022fd6ab26ce0b9	prod-nginx-sg	vpc-0da180d315adb9e93	Nginx reverse proxy
<input type="checkbox"/>	dev-sg	sg-0ee9dce4b4e7c9999	default	vpc-01825459d58b35672	default VPC security
<input type="checkbox"/>	prod-backend-sg	sg-07a0c33cf9396db51	prod-backend-sg	vpc-0da180d315adb9e93	Backend web server
<input type="checkbox"/>	-	sg-08511e2a278e0ffdf34	default	vpc-01b7d554e460c902d	default VPC security
<input type="checkbox"/>	-	sg-0a6a674172092bd49	default	vpc-0da180d315adb9e93	default VPC security

The screenshot shows the AWS Security Groups console for a security group named 'prod-nginx-sg'. Key details include:

- Security group name:** prod-nginx-sg
- Security group ID:** sg-01022fd6ab26ce0b9
- Description:** Nginx reverse proxy / load balancer security group
- VPC ID:** vpc-0da180d315adb9e93
- Owner:** 05073780689
- Inbound rules count:** 3 Permission entries
- Outbound rules count:** 1 Permission entry

The 'Inbound rules' tab is selected, showing three rules:

Security group rule ID	IP version	Type	Protocol	Port range	Source
sgr-0f5df2eb01b35cf6d	IPv4	SSH	TCP	22	154.192.18.10/32
sgr-08ca1ef36505e307d	IPv4	HTTP	TCP	80	0.0.0.0/0
sgr-09eeee7327aa10cf5	IPv4	HTTPS	TCP	443	0.0.0.0/0

1.5 Locals Configuration (5 marks)

Create locals.tf with:

Dynamic IP detection for my_ip

Resource naming conventions

Common tags

Backend server configurations

Example:

```
locals {
```

```
    my_ip = "${chomp(data.http.my_ip.response_body)}/32"
```

```
    common_tags = {
```

```
        Environment = var.env_prefix
```

```
        Project     = "Assignment-2"
```

```
        ManagedBy   = "Terraform"
```

```
}
```

```
backend_servers = [
```

```
{
```

```
    name      = "web-1"
```

```
    suffix    = "1"
```

```
    script_path = "./scripts/apache-setup.sh"
```

```
,
```



```
    name      = "web-2"
```

```
    suffix    = "2"
```

```
    script_path = "./scripts/apache-setup.sh"
```

```
,
```



```
    name      = "web-3"
```

```

    suffix      = "3"
    script_path = "./scripts/apache-setup.sh"
}
]
}
data "http" "my_ip" {
  url = "https://icanhazip.com"
}

```

Tasks:

- Implement dynamic IP detection
- Define common tags
- Define backend server list
- Add any other reusable local values

Deliverables:

Screenshot: assignment_part1_locals_tf.png

```

# Fetch your public IP address dynamically from ianahazip.com
data "http" "my_ip" {
  url = "https://icanhazip.com"
}

locals {
  # Dynamic IP with /32 CIDR for security group
  my_ip = "${chomp(data.http.my_ip.response_body)}/32"

  # Common tags to apply to all resources
  common_tags = {
    Environment = var.env_prefix
    Project     = "Assignment-2"
    ManagedBy   = "Terraform"
  }
}

# List of backend servers with names, suffix, and scripts
backend_servers = [
  {
    name      = "web-1"
    suffix    = "1"
    script_path = "./scripts/apache-setup.sh"
  },
  {
    name      = "web-2"
    suffix    = "2"
    script_path = "./scripts/apache-setup.sh"
  },
  {
    name      = "web-3"
  }
]

```

Part 2:

Webserver Module (15 marks)

2.1 Module Design (10 marks)

Create a reusable webserver module in modules/webserver/

Module Requirements:

Variables:

env_prefix

instance_name

instance_type

availability_zone

vpc_id

subnet_id

security_group_id

public_key

script_path

instance_suffix

common_tags

Resources to Create:

AWS Key Pair (unique per instance using suffix)

EC2 Instance with:

Specified AMI (Amazon Linux 2023)

Public IP enabled

User data from script file

Proper tags

Outputs:

instance_id

public_ip

private_ip

Tasks:

Create variables.tf with all required variables

Create main.tf with key pair and instance resources

Create outputs.tf with all required outputs

Ensure module is reusable for both Nginx and backend servers

Deliverables:

Screenshot: assignment_part2_webserver_module_variables.png

g

```
variable "subnet_id" {
  type     = string
  description = "Subnet ID to launch instance in"
}

variable "security_group_id" {
  type     = string
  description = "Security Group ID to associate with instance"
}

variable "public_key" {
  type     = string
  description = "Path to SSH public key file"
}

variable "script_path" {
  type     = string
  description = "Path to user data script"
}

variable "instance_suffix" {
  type     = string
  description = "Unique suffix for the instance"
}

variable "common_tags" {
  type     = map(string)
  description = "Common tags applied to resources"
}
```

Screenshot: assignment_part2_webserver_module_main.png

```

# Read public key content
resource "tls_private_key" "example" {
  algorithm = "RSA"
  rsa_bits  = 4096
}

resource "aws_key_pair" "web_key" {
  key_name   = "${var.env_prefix}-${var.instance_name}-${var.instance_suffix}-key"
  public_key = file(var.public_key)

  tags = merge(var.common_tags, {
    Name = "${var.env_prefix}-${var.instance_name}-${var.instance_suffix}-key"
  })
}

resource "aws_instance" "web_instance" {
  ami           = "ami-05524d6658fcf35b6" # Amazon Linux 2023 (update if needed)
  instance_type = var.instance_type
  subnet_id     = var.subnet_id
  vpc_security_group_ids = [var.security_group_id]
  key_name      = aws_key_pair.web_key.key_name
  associate_public_ip_address = true
  availability_zone = var.availability_zone

  user_data = file(var.script_path)

  tags = merge(var.common_tags, {
    Name = "${var.env_prefix}-${var.instance_name}-${var.instance_suffix}"
  })
}

```

Screenshot: assignment_part2_webserver_module_outputs.png

```

output "instance_id" {
  description = "ID of the EC2 instance"
  value       = aws_instance.web_instance.id
}

output "public_ip" {
  description = "Public IP address of the EC2 instance"
  value       = aws_instance.web_instance.public_ip
}

output "private_ip" {
  description = "Private IP address of the EC2 instance"
  value       = aws_instance.web_instance.private_ip
}

```

2.2 Module Usage (5 marks)

In root main.tf, instantiate the webserver module for:

One Nginx server (using nginx-setup.sh)

Three backend servers (web-1, web-2, web-3 using apache-setup.sh)

Tasks:

- Create Nginx server module instance
- Create backend server module instances using for_each
- Ensure proper security group assignment
- Pass all required variables

Deliverables:

Screenshot: assignment_part2_main_tf_modules.png

```

module "nginx_server" {
  source          = "./modules/webserver"
  env_prefix      = var.env_prefix
  instance_name   = "nginx-proxy"
  instance_type   = var.instance_type
  availability_zone = var.availability_zone
  vpc_id          = module.networking.vpc_id
  subnet_id       = module.networking.subnet_id
  security_group_id = module.security.nginx_sg_id
  public_key       = var.public_key
  script_path      = "./scripts/nginx-setup.sh"
  instance_suffix  = "nginx"
  common_tags     = local.common_tags
}

# Backend Servers (Apache Web Servers)
module "backend_servers" {
  for_each = { for idx, server in local.backend_servers : server.name => server }

  source          = "./modules/webserver"
  env_prefix      = var.env_prefix
  instance_name   = each.value.name
  instance_type   = var.instance_type
  availability_zone = var.availability_zone
  vpc_id          = module.networking.vpc_id
  subnet_id       = module.networking.subnet_id
  security_group_id = module.security.backend_sg_id
  public_key       = var.public_key
  script_path      = each.value.script_path
  instance_suffix  = each.value.suffix
  common_tags     = local.common_tags
}

```

Part 3: Server Configuration Scripts (20 marks)

3.1 Apache Backend Server Script (10 marks)

Create scripts/apache-setup.sh that:

Updates system packages

Installs Apache HTTP Server

Starts and enables Apache service

Creates custom HTML page displaying:

Server name/hostname

Private IP address

Public IP address

Public DNS hostname

Custom message indicating which backend server (web-1, web-2, or web-3)

Timestamp of deployment

Server status (Primary/Backup)

Tasks:

- Create the script with all required features
- Ensure it uses IMDSv2 for metadata
- Create visually appealing HTML output
- Make script executable
- Test script independently

Deliverables:

Screenshot: assignment_part3_apache_script.png

```

#!/bin/bash
set -e

# Update system
yum update -y

# Install Apache
yum install httpd -y

# Start and enable Apache
systemctl start httpd
systemctl enable httpd

# Get metadata token (IMDSv2)
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" \
-H "X-aws-ec2-metadata-token-ttl-seconds: 21600")

# Get instance metadata
PRIVATE_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/local-ipv4)
PUBLIC_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-ipv4)
PUBLIC_DNS=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-hostname)
INSTANCE_ID=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/instance-id)

# Set hostname
hostnamectl set-hostname myapp-webserver

```

Screenshot: assignment_part3_backend_webpage.png (browser showing backend server page)



3.2 Nginx Server Setup Script (10 marks)

Create scripts/nginx-setup.sh that:
Updates system packages
Installs Nginx
Generates self-signed SSL certificate
Configures Nginx with:
HTTPS on port 443
HTTP to HTTPS redirect
Upstream backend servers (web-1, web-2, web-3)
Load balancing configuration
Caching configuration

Proper logging

Security headers

Tasks:

- Create script with SSL certificate generation
- Configure upstream with placeholder IPs
- Implement caching
- Add security headers
- Configure HTTP to HTTPS redirect
- Add health check endpoint

Deliverables:

Screenshot: assignment_part3_nginx_script.png

```
#!/bin/bash
set -e

# Update and install Nginx and SSL
yum update -y
yum install -y nginx openssl
systemctl start nginx
systemctl enable nginx

# Create SSL directories
mkdir -p /etc/ssl/private
mkdir -p /etc/ssl/certs

# Get metadata token (IMDSv2)
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" \
-H "X-aws-ec2-metadata-token-ttl-seconds: 21600")

# Get public IP
PUBLIC_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-ipv4)

# Generate self-signed certificate
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \
-keyout /etc/ssl/private/selfsigned.key \
-out /etc/ssl/certs/selfsigned.crt \
-subj "/CN=$PUBLIC_IP" \
-adext "subjectAltName=IP:$PUBLIC_IP" \
-adext "basicConstraints=CA:FALSE" \
-adext "keyUsage=digitalSignature,keyEncipherment" \
-adext "extendedKeyUsage=serverAuth"
```

Screenshot: assignment_part3_nginx_default_page.png (before backend configuration)

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

Part 4:

Infrastructure Deployment (15 marks)

4.1 Initial Deployment (5 marks)

Deploy the infrastructure using Terraform.

Tasks:

- Generate SSH key pair if not exists
- Initialize Terraform (terraform init)
- Validate configuration (terraform validate)
- Plan deployment (terraform plan)
- Apply configuration (terraform apply -auto-approve)

Deliverables:

Screenshot: assignment_part4_ssh_keygen.png

```
PS C:\Users\user\OneDrive\Desktop\Assignment2> ls ~/ssh/
```

```
Directory: C:\Users\user\.ssh
```

Mode	LastWriteTime	Length	Name
----	-----	-----	-----
-a---	9/27/2025 4:16 PM	419	id_ed25519
-a---	9/27/2025 4:16 PM	110	id_ed25519.pub

Screenshot: assignment_part4_terraform_init.png

```
PS C:\Users\user\OneDrive\Desktop\Assignment2> terraform init
```

```
Initializing the backend...  
Initializing modules...
```

```
Initializing provider plugins...
```

```
- Reusing previous version of hashicorp/aws from the dependency lock file  
- Reusing previous version of hashicorp/tls from the dependency lock file  
- Reusing previous version of hashicorp/http from the dependency lock file  
- Using previously-installed hashicorp/aws v6.27.0  
- Using previously-installed hashicorp/tls v4.1.0  
- Using previously-installed hashicorp/http v3.5.0
```

```
Terraform has been successfully initialized!
```

```
You may now begin working with Terraform. Try running "terraform plan" to see  
any changes that are required for your infrastructure. All Terraform commands  
should now work.
```

```
If you ever set or change modules or backend configuration for Terraform,  
rerun this command to reinitialize your working directory. If you forget, other  
commands will detect it and remind you to do so if necessary.  
PS C:\Users\user\OneDrive\Desktop\Assignment2> |
```

Screenshot: assignment_part4_terraform_validate.png

```
PS C:\Users\user\OneDrive\Desktop\Assignment2> terraform validate  
Success! The configuration is valid.
```

```
PS C:\Users\user\OneDrive\Desktop\Assignment2> |
```

Screenshot: assignment_part4_terraform_plan.png

```

PS C:\Users\user\OneDrive\Desktop\Assignment2> terraform plan
data.http_my_ip: Reading...
module.nginx_server.tls_private_key.example: Refreshing state... [id=8937bb19f82ddbb73f5037f3a6cccaae95bdca23]
module.backend_servers["web-1"].tls_private_key.example: Refreshing state... [id=4d45ed3887e2031ae86e4cc555f36f7f897e9a1]
module.backend_servers["web-3"].tls_private_key.example: Refreshing state... [id=2a118e6f78dd3135ba929337acfcb4b258801fe8]
module.backend_servers["web-2"].tls_private_key.example: Refreshing state... [id=cd14656e009651923ad7736a0ef6763c1b812be9]
data.http_my_ip: Read complete after [id=https://icanhazip.com]

module.backend_servers["web-2"].aws_key_pair.web_key: Refreshing state... [id=prod-web-2-2-key]
module.backend_servers["web-3"].aws_key_pair.web_key: Refreshing state... [id=prod-web-3-3-key]
module.backend_servers["web-1"].aws_key_pair.web_key: Refreshing state... [id=prod-web-1-1-key]
module.nginx_server.aws_key_pair.web_key: Refreshing state... [id=prod-nginx-proxy-nginx-key]
module.networking.aws_vpc.main: Refreshing state... [id=vc-0da189d315adb9e93]
module.networking.aws_internet_gateway.main: Refreshing state... [id=igw-0b2943b282febfb7]
module.networking.aws_subnet.main: Refreshing state... [id=subnet-082ab45a8debe2b2]
module.security.aws_security_group.nginx: Refreshing state... [id=sg-01022fd6ab26cebb9]
module.networking.aws_route_table.main: Refreshing state... [id=rtb-09cf58ad16177c493]
module.security.aws_security_group.backend: Refreshing state... [id=sg-07a0c33cf9396db51]
module.nginx_server.aws_instance.web_instance: Refreshing state... [id=i-02ec91f70c9f1227b]
module.networking.aws_route_table_association.main: Refreshing state... [id=rtbassoc-0818910187e73934d]
module.backend_servers["web-2"].aws_instance.web_instance: Refreshing state... [id=i-03ed16575f0c734ae]
module.backend_servers["web-3"].aws_instance.web_instance: Refreshing state... [id=i-072abf049f195f04b]
module.backend_servers["web-1"].aws_instance.web_instance: Refreshing state... [id=i-0048cc9baa2df4137]

```

Note: Objects have changed outside of Terraform

Terraform detected the following changes made outside of Terraform since the last "terraform apply" which may have affected this plan:

```

# module.backend_servers["web-1"].aws_instance.web_instance has changed
~ resource "aws_instance" "web_instance" {
  id          = "i-0048cc9baa2df4137"
  ~ public_ip   = "51.112.46.85" -> "158.252.114.208"
  tags        = {
    "Environment" = "prod"
    "ManagedBy"   = "Terraform"
    "Name"        = "prod-web-1-1"
    "Project"     = "Assignment-2"
  }
}

```

Screenshot: assignment_part4_terraform_apply.png (showing all 4 instances created)

Apply complete! Resources: 0 added, 2 changed, 0 destroyed.

Outputs:

```

backend_server_instance_ids = [
  "web-1" = "i-0048cc9baa2df4137"
  "web-2" = "i-03ed16575f0c734ae"
  "web-3" = "i-072abf049f195f04b"
]
backend_server_private_ips = {
  "web-1" = "10.0.10.178"
  "web-2" = "10.0.10.168"
  "web-3" = "10.0.10.237"
}
backend_server_public_ips = {
  "web-1" = "158.252.114.208"
  "web-2" = "51.112.58.134"
  "web-3" = "3.28.130.127"
}

```

4.2 Output Configuration (5 marks)

Tasks:

Create all required outputs

Add helpful descriptions

Include configuration guide

Display outputs after apply

Commands:

terraform output

terraform output -json > outputs.json

Deliverables:

Screenshot: assignment_part4_terraform_output.png

```
=====
DEPLOYMENT SUCCESSFUL!
=====

Next Steps:
1. SSH into Nginx server: ssh ec2-user@158.252.32.185
2. Edit Nginx config: sudo vim /etc/nginx/nginx.conf
3. Update backend IPs in upstream block:
   - BACKEND_IP_1: 10.0.10.178
   - BACKEND_IP_2: 10.0.10.168
   - BACKEND_IP_3: 10.0.10.237
4. Restart Nginx: sudo systemctl restart nginx
5. Test: https://158.252.32.185

Backend Servers:
- web-1: 158.252.114.208 (private: 10.0.10.178)
  - web-2: 51.112.50.134 (private: 10.0.10.168)
  - web-3: 3.28.130.127 (private: 10.0.10.237)

=====
EOT
nginx_instance_id = "i-02ec91f70c9f1227b"
nginx_public_ip = "158.252.32.185"
subnet_id = "subnet-082ab45a8debe2bb2"
vpc_id = "vpc-0da180d315adb9e93"
```

Screenshot: assignment_part4_outputs_json.png

```
{
  "web-3": {
    "instance_id": "i-072abf049f195f04b",
    "private_ip": "10.0.10.237",
    "public_ip": "3.28.130.127"
  },
  "configuration_guide": {
    "sensitive": false,
    "type": "string",
    "value": "=====\\r\\nDEPLOYMENT SUCCESSFUL!\\r\\n=====\\r\\n====\\r\\nNext Steps:\\r\\n1. Ssh into Nginx server: ssh ec2-user@158.252.32.185\\r\\n2. Edit Nginx config: sudo vim /etc/nginx/nginx.conf\\r\\n3. Update backend IPs in upstream block:\\r\\nBACKEND_IP_1: 10.0.10.178\\r\\n - BACKEND_IP_2: 10.0.10.168\\r\\n - BACKEND_IP_3: 10.0.10.237\\r\\n4. Restart Nginx: sudo systemctl restart nginx\\r\\n5. Test: https://158.252.32.185\\r\\nBackend Servers:\\r\\n- web-1: 158.252.114.208 (private: 10.0.10.178)\\r\\n- web-2: 51.112.50.134 (private: 10.0.10.168)\\r\\n- web-3: 3.28.130.127 (private: 10.0.10.237)\\r\\n\\r\\n=====\\r\\n"
  },
  "nginx_instance_id": {
    "sensitive": false,
    "type": "string",
    "value": "i-02ec91f70c9f1227b"
  },
  "nginx_public_ip": {
    "sensitive": false,
    "type": "string",
    "value": "158.252.32.185"
  },
  "subnet_id": {
    "sensitive": false,
    "type": "string",
    "value": "subnet-082ab45a8debe2bb2"
  },
  "vpc_id": {
    "sensitive": false,
    "type": "string",
    "value": "vpc-0da180d315adb9e93"
  }
}
```

4.3 AWS Console Verification (5 marks)

Verify all resources in AWS Console.

Tasks:

- Verify VPC created
- Verify Subnet created
- Verify Internet Gateway attached
- Verify Route Table configured
- Verify Security Groups created with correct rules
- Verify all 4 EC2 instances running
- Verify Key Pairs created

Deliverables:

Screenshot: assignment_part4_aws_vpc.png (VPC console)

Your VPCs

VPCs		VPC encryption controls		Actions		Create VPC
Your VPCs (3) Info				Last updated	Actions	Create VPC
		Find VPCs by attribute or tag		Last updated	Actions	Create VPC
Name	VPC ID	State	Encryption c...	Encryption control ...	Block Public...	IPv4
dev-vpc	vpc-01825459d38b35672	Available	-	-	Off	10.0.0.0/16
prod-vpc	vpc-0da180d315adb9e93	Available	-	-	Off	10.0.0.0/16
-	vpc-01b7d554ea60c902d	Available	-	-	Off	172.31.0.0/20

Select a VPC above

Screenshot: assignment_part4_aws_subnet.png (Subnet details)

Subnets (5) Info

Subnets (5) Info		Actions		Create subnet
		Last updated	Actions	Create subnet
		Find subnets by attribute or tag	Last updated	Actions
Name	Subnet ID	State	VPC	Block Public... IPv4 CIDR
-	subnet-0b11cd485731a00f7	Available	vpc-01b7d554ea60c902d	Off 172.31.32.0/24
prod-subnet	subnet-082ab45a8dbebe2bb2	Available	vpc-0da180d315adb9e93 prod-vpc	Off 10.0.10.0/24
-	subnet-09a23e3b14a74df6e	Available	vpc-01b7d554ea60c902d	Off 172.31.0.0/20
-	subnet-06e5d96b7c34d750f	Available	vpc-01b7d554ea60c902d	Off 172.31.16.0/24

Screenshot: assignment_part4_aws_security_groups.png (Both security groups)

sg-07a0c33cf9396db51 - prod-backend-sg

Inbound rules (2)

Inbound rules (2)		Manage tags		Edit inbound rules	
		Last updated	Actions	Create rule	Copy
Name	Security group rule ID	IP version	Type	Protocol	Port range
-	sgr-0adc1bc183bb4995f	IPv4	SSH	TCP	22
-	sgr-03e4f2137c0277e56	-	HTTP	TCP	80

sg-01022fd6ab26ce0b9 - prod-nginx-sg

Inbound rules (3)

Inbound rules (3)		Manage tags		Edit inbound rules	
		Last updated	Actions	Create rule	Copy
Name	Security group rule ID	IP version	Type	Protocol	Port range
-	sgr-0f5df2eb01b35cf6d	IPv4	SSH	TCP	22
-	sgr-08ca1ef36505e307d	IPv4	HTTP	TCP	80
-	sgr-09eeee7327aa10cf5	IPv4	HTTPS	TCP	443

Screenshot: assignment_part4_aws_instances.png (All 4 running instances)

Instance ID	Instance Type	Status	Health	Placement	Actions
i-0048cc9baa2df4137	t3.micro	Running	3/3 checks passed	me-central-1a	
i-02ec91f70c9f1227b	t3.micro	Running	3/3 checks passed	me-central-1a	
i-03ed16575f0c734ae	t3.micro	Running	3/3 checks passed	me-central-1a	
i-0c38fe3027793a0b7	t3.micro	Running	3/3 checks passed	me-central-1a	
i-072abf049f195f04b	t3.micro	Running	3/3 checks passed	me-central-1a	

Part 5:

Nginx Configuration & Testing (25 marks)

5.1 Update Nginx Backend Configuration (5 marks)

SSH into the Nginx server and update the configuration with actual backend IPs.

Tasks:

- SSH into Nginx server
- Edit /etc/nginx/nginx.conf

- Replace placeholder IPs with actual private IPs of backend servers
- Test Nginx configuration
- Restart Nginx service

Deliverables:

Screenshot: assignment_part5_ssh_nginx.png (SSH session)

```
PS C:\Users\user\OneDrive\Desktop\Assignment2> ssh ec2-user@158.252.32.185
The authenticity of host '158.252.32.185' (158.252.32.185) can't be established.
ED25519 key fingerprint is SHA256:hLNimd3WgQz9dd40ZIZsUphj2BZwoHPPEPLBFCYxoA.
This host key is known by the following other names/addresses:
  C:\Users\user/.ssh/known_hosts:9: 51.112.252.77
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '158.252.32.185' (ED25519) to the list of known hosts.

  _\_
  ~\_\_ #####_      Amazon Linux 2023
  ~\_\_ \_\#\#\#`_
  ~\_\_ \#\#\#
  ~\_\_ \#/` --- https://aws.amazon.com/linux/amazon-linux-2023
  ~\_\_ \V~` '--->
  ~\_\_ /` /
  ~\_\_ /` /
  ~\_\_ /` /
  ~\_\_ /` /
  ~\_\_ /` /
>Last login: Fri Dec 26 19:14:46 2025 from 154.192.19.36
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: assignment_part5_nginx_conf_updated.png (updated upstream block)

```
upstream backend_servers {
    least_conn;
    server 10.0.10.178:80 max_fails=3 fail_timeout=30s; # web-1
    server 10.0.10.168:80 max_fails=3 fail_timeout=30s; # web-2
    server 10.0.10.237:80; # web-3 (backup)
}
```

Screenshot: assignment_part5_nginx_test.png (nginx -t output)

```
[ec2-user@ip-10-0-10-242 ~]$ sudo nginx -t
nginx: [warn] could not build optimal types_hash, you should increase either types_hash_max_size: 1024 or types_hash_bucket_size
nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
nginx: configuration file /etc/nginx/nginx.conf test is successful
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: assignment_part5_nginx_restart.png (restart and status)

```
nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
nginx: configuration file /etc/nginx/nginx.conf test is successful
[ec2-user@ip-10-0-10-242 ~]$ sudo systemctl start nginx
[ec2-user@ip-10-0-10-242 ~]$ sudo systemctl status nginx
● nginx.service - The nginx HTTP and reverse proxy server
   Loaded: loaded (/usr/lib/systemd/system/nginx.service; enabled; preset: disabled)
   Active: active (running) since Fri, 2025-12-26 20:33:02 UTC; 45ms ago
     Process: 3985 ExecStart=/usr/bin/nginx -f /run/nginx.pid (code=exited, status=0/SUCCESS)
    Process: 3986 ExecStartPre=/usr/sbin/nginx -t (code=exited, status=0/SUCCESS)
    Process: 3987 ExecStart=/usr/sbin/nginx (code=exited, status=0/SUCCESS)
   Main PID: 3988 (nginx)
      Tasks: 3 (limit: 1067)
        Memory: 3.3M
          CPU: 45ms
         CGroup: /system.slice/nginx.service
             ├─3988 "nginx: master process /usr/sbin/nginx"
             ├─3989 "nginx: worker process"
             └─3990 "nginx: worker process"

Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal systemd[1]: Starting nginx.service - The nginx HTTP and reverse proxy server...
Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal nginx[3986]: nginx: [warn] could not build optimal types_hash, you should increase either type
Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal nginx[3986]: nginx: the configuration file /etc/nginx/nginx.conf syntax is ok
Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal nginx[3986]: nginx: configuration file /etc/nginx/nginx.conf test is successful
Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal nginx[3987]: nginx: [warn] could not build optimal types_hash, you should increase either type
Dec 26 20:33:02 ip-10-0-10-242.me-central-1.compute.internal systemd[1]: Started nginx.service - The nginx HTTP and reverse proxy server.

[ec2-user@ip-10-0-10-242 ~]$ |
```

5.2 Test Load Balancing (5 marks)

Test that Nginx is properly load balancing between web-1 and web-2.

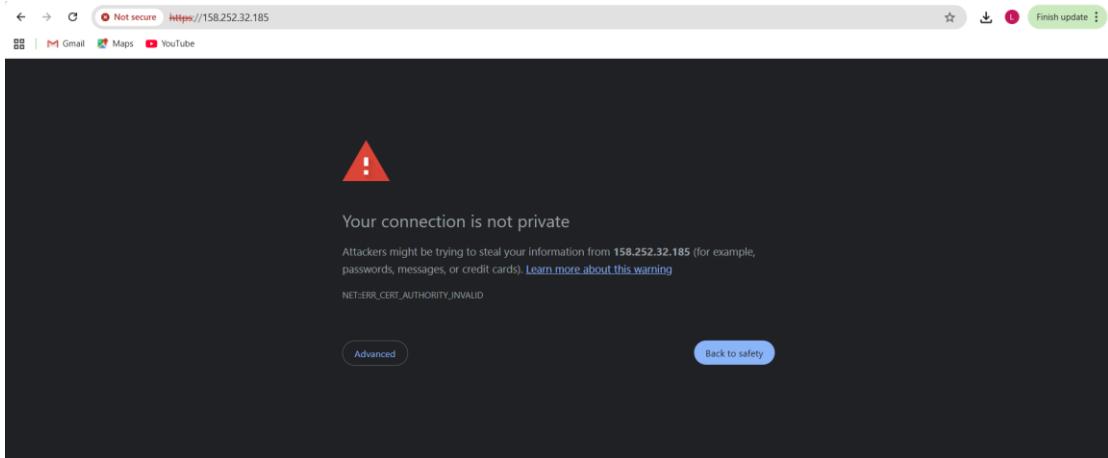
Tasks:

- Open browser to <https://<nginx-public-ip>>
- Accept the security warning for self-signed certificate
- Reload page multiple times (at least 10 times)
- Verify traffic alternates between web-1 and web-2
- Verify web-3 is NOT serving traffic (it's backup only)

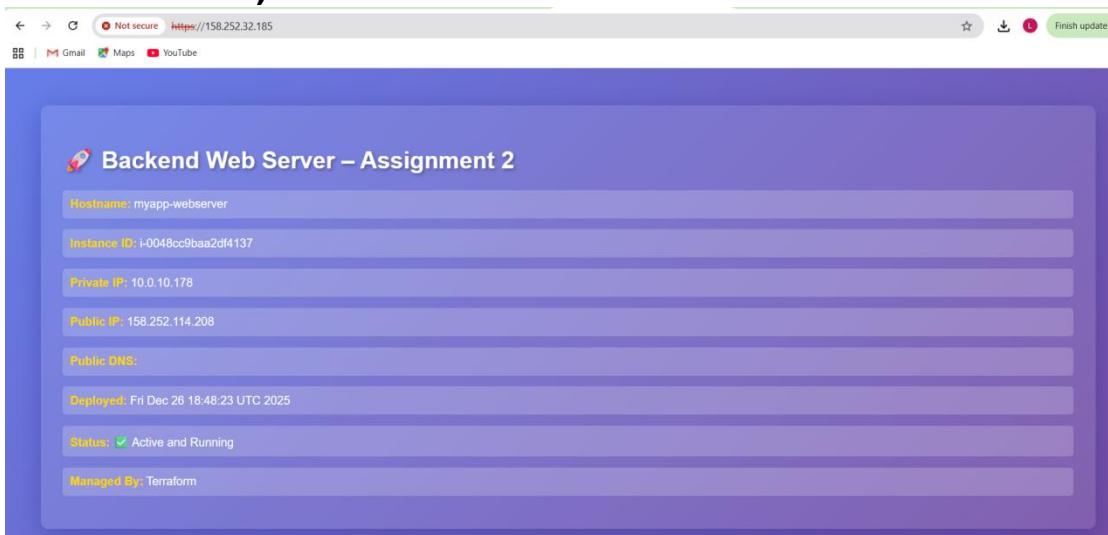
- Document the load balancing pattern

Deliverables:

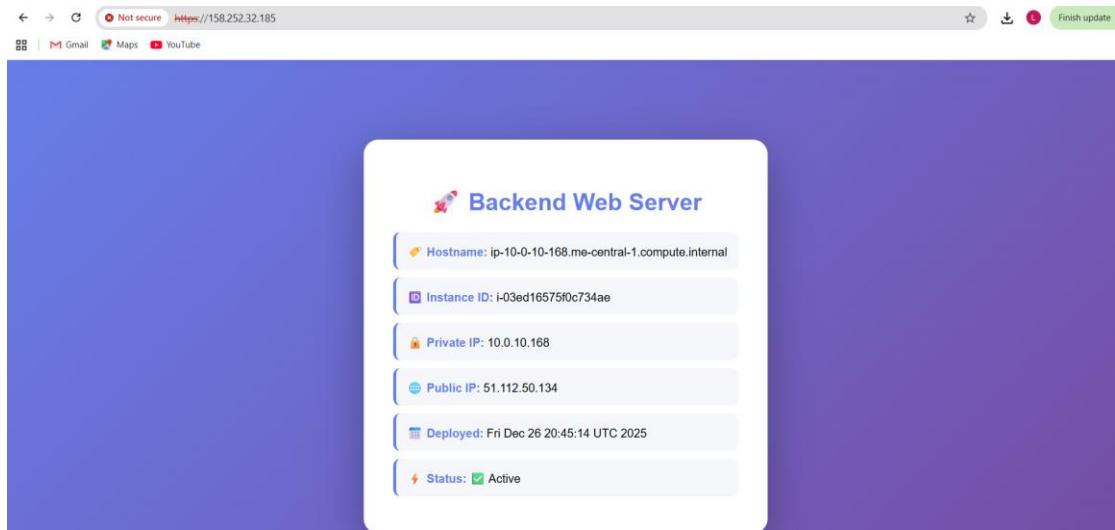
Screenshot: assignment_part5_ssl_warning.png (browser security warning)



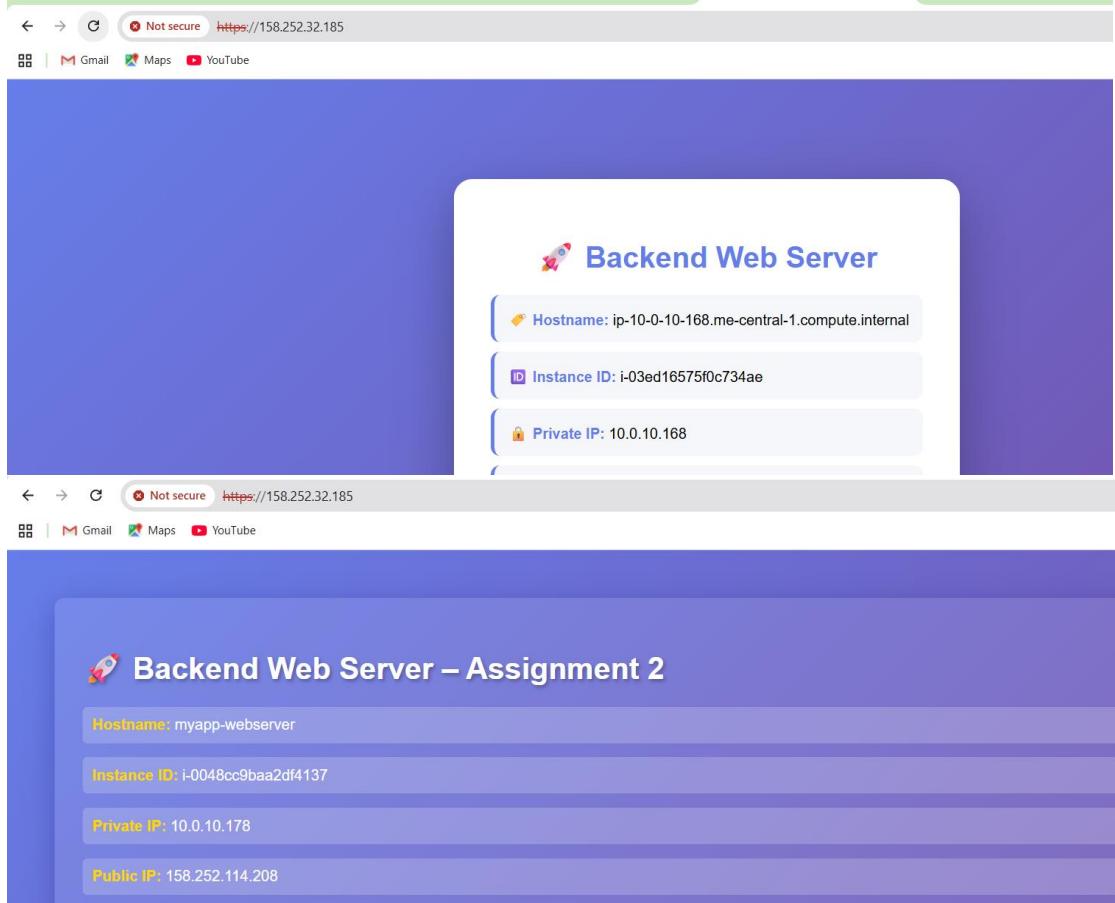
Screenshot: assignment_part5_web1_response.png (showing web-1 content)



Screenshot: assignment_part5_web2_response.png (showing web-2 content)



**Screenshot: assignment_part5_load_balancing_demo.png
(multiple reloads showing alternation)**



5.3 Test Cache Functionality (5 marks) Verify that Nginx caching is working correctly.

Tasks:

- Open browser developer tools (F12)
- Navigate to Network tab
- Clear browser cache
- Load `https://<nginx-public-ip>`

- Check response headers for X-Cache-Status: MISS (first request)
- Reload page
- Check response headers for X-Cache-Status: HIT (cached request)
- Verify cache directory on Nginx server

Deliverables:

Screenshot: assignment_part5_cache_miss.png (first request - MISS)

```
[ec2-user@ip-10-0-10-242 ~]$ curl -k -I "https://localhost/" | grep -i "x-cache-status"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left Speed
0  1622     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0
x-cache-status: MISS
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: assignment_part5_cache_hit.png (second request - HIT)

```
[ec2-user@ip-10-0-10-242 ~]$ curl -k -I "https://localhost/" | grep -i "x-cache-status"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left Speed
0  1622     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0
x-cache-status: HIT
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: assignment_part5_cache_directory.png (cache folder contents)

```
[ec2-user@ip-10-0-10-242 ~]$ ls -la /var/cache/nginx/
total 0
drwxr-xr-x. 3 nginx nginx 19 Dec 26 21:18 .
drwxr-xr-x. 10 root root 114 Dec 26 21:18 ..
drwxr-xr-x. 5 nginx nginx 33 Dec 26 21:37 cache
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: assignment_part5_access_log_cache.png (access log showing cache status)

```
154.192.19.36 - - [26/Dec/2025:21:34:30 +0000] "GET / HTTP/2.0" 200 703 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.168:80
154.192.19.36 - - [26/Dec/2025:21:34:32 +0000] "GET / HTTP/2.0" 200 647 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80
::1 - - [26/Dec/2025:21:37:01 +0000] "HEAD / HTTP/2.0" 200 0 "-" "curl/8.11.1" "-" Cache: MISS Backend: 10.0.10.168:80
::1 - - [26/Dec/2025:21:37:42 +0000] "HEAD / HTTP/2.0" 200 0 "-" "curl/8.11.1" "-" Cache: HIT Backend: -
[ec2-user@ip-10-0-10-242 ~]$ |
```

5.4 Test High Availability (Backup Server) (5 marks)

Test the backup server functionality by simulating primary server failure.

Tasks:

- SSH into web-1 and stop Apache
- Reload Nginx page - should show web-2 only
- SSH into web-2 and stop Apache
- Reload Nginx page - should now show web-3 (backup activated)
- Restart web-1 and web-2
- Verify traffic returns to web-1 and web-2

Deliverables:

Screenshot: assignment_part5_web1_stopped.png (web-1 Apache stopped)

```

>Last login: Fri Dec 26 21:04:04 2025 from 154.192.19.36
[ec2-user@myapp-webserver ~]$ sudo systemctl stop httpd
[ec2-user@myapp-webserver ~]$ sudo systemctl status httpd
○ httpd.service - The Apache HTTP Server
    Loaded: loaded (/usr/lib/systemd/system/httpd.service; disabled; preset: disabled)
      Active: inactive (dead)
        Docs: man:httpd.service(8)

Dec 26 21:04:24 myapp-webserver systemd[1]: httpd.service: Failed with result 'exit-code'.
Dec 26 21:04:24 myapp-webserver systemd[1]: Failed to start httpd.service - The Apache HTTP Server.
Dec 26 21:05:06 myapp-webserver systemd[1]: Starting httpd.service - The Apache HTTP Server...
Dec 26 21:05:06 myapp-webserver httpd[36421]: AH00558: httpd: Could not reliably determine the server's fully qualified name, using 154.192.19.36 for Port 80
Dec 26 21:05:06 myapp-webserver systemd[1]: Started httpd.service - The Apache HTTP Server.
Dec 26 21:05:06 myapp-webserver httpd[36421]: Server configured, listening on: port 80
Dec 26 21:42:23 myapp-webserver systemd[1]: Stopping httpd.service - The Apache HTTP Server...
Dec 26 21:42:24 myapp-webserver systemd[1]: httpd.service: Deactivated successfully.
Dec 26 21:42:24 myapp-webserver systemd[1]: Stopped httpd.service - The Apache HTTP Server.
Dec 26 21:42:24 myapp-webserver systemd[1]: httpd.service: Consumed 1.932s CPU time.

```

Screenshot: assignment_part5_web2_stopped.png (web-2 Apache stopped)

```

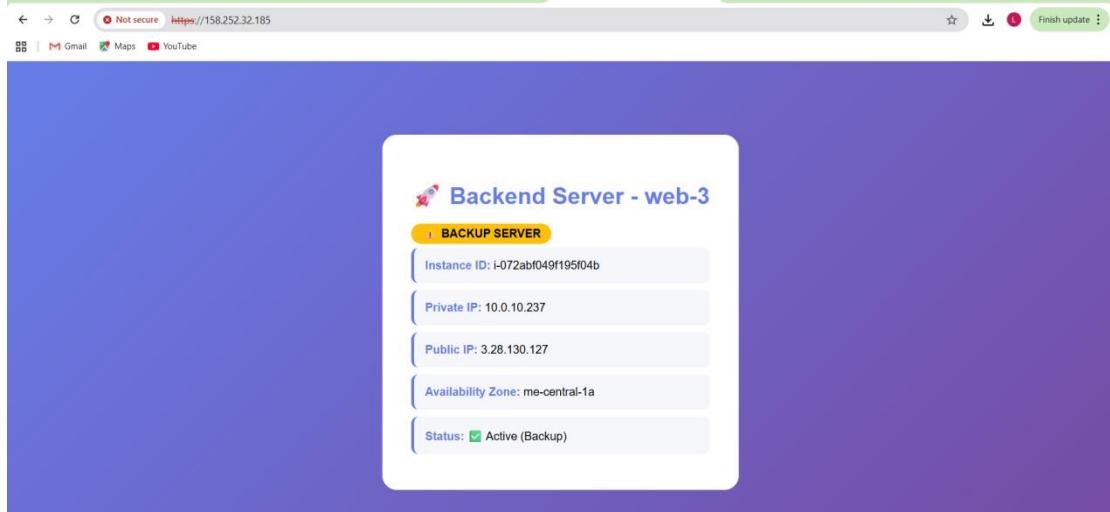
Last login: Fri Dec 26 20:43:46 2025 from 154.192.19.36
[ec2-user@ip-10-0-10-168 ~]$ sudo systemctl stop httpd

# Verify it's stopped
sudo systemctl status httpd
○ httpd.service - The Apache HTTP Server
    Loaded: loaded (/usr/lib/systemd/system/httpd.service; enabled; preset: disabled)
      Active: inactive (dead) since Fri 2025-12-26 21:44:32 UTC; 114ms ago
        Duration: 59min 45.566s
          Docs: man:httpd.service(8)
        Process: 9972 ExecStart=/usr/sbin/httpd $OPTIONS -DFOREGROUND (code=exited, status=0/SUCCESS)
      Main PID: 9972 (code=exited, status=0/SUCCESS)
        Status: "Total requests: 16; Idle/Busy workers 100/0; Requests/sec: 0.00447; Bytes served/sec: 10 B/sec"
          CPU: 3.462s

Dec 26 20:44:45 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Starting httpd.service - The Apache HTTP Server...
Dec 26 20:44:45 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Started httpd.service - The Apache HTTP Server.
Dec 26 20:44:45 ip-10-0-10-168.me-central-1.compute.internal httpd[9972]: Server configured, listening on: port 80
Dec 26 21:44:31 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Stopping httpd.service - The Apache HTTP Server...
Dec 26 21:44:32 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: httpd.service: Deactivated successfully.
Dec 26 21:44:32 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Stopped httpd.service - The Apache HTTP Server.
Dec 26 21:44:32 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: httpd.service: Consumed 3.462s CPU time.
[ec2-user@ip-10-0-10-168 ~]$ ^C
[ec2-user@ip-10-0-10-168 ~]$ 

```

Screenshot: assignment_part5_backup_activated.png (web-3 serving traffic)



Screenshot: assignment_part5_nginx_error_log.png (error log showing backend failures)

```

2025/12/26 21:45:15 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to upstream, client: 154.192.19.36
r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream, client: 154.192.19.36
equest: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to upstream, client: 154.192.19.36
r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream, client: 154.192.19.36
equest: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to upstream, client: 154.192.19.36
r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream, client: 154.192.19.36
equest: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to upstream, client: 154.192.19.36
r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream, client: 154.192.19.36
equest: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"

```

Screenshot: assignment_part5_services_restored.png (all services back online)

```

Last login: Fri Dec 26 21:48:03 2025 from 154.192.19.36
[ec2-user@myapp-websvserver ~]$ sudo systemctl enable httpd
Created symlink /etc/systemd/system/multi-user.target.wants/httpd.service → /usr/lib/systemd/system/httpd.service.
[ec2-user@myapp-websvserver ~]$ sudo systemctl start httpd
[ec2-user@myapp-websvserver ~]$ sudo systemctl status httpd
● httpd.service - The Apache HTTP Server
   Loaded: loaded (/usr/lib/systemd/system/httpd.service; enabled; preset: disabled)
   Active: active (running) since Fri 2025-12-26 21:50:54 UTC; 6s ago
     Docs: man:httpd.service(8)
   Main PID: 38813 (httpd)
      Status: "Started, listening on: port 80"
         Tasks: 177 (limit: 1067)
        Memory: 13.4M
       CPU: 64ms
      CGroup: /system.slice/httpd.service
              ├─38813 /usr/sbin/httpd -DFOREGROUND
              ├─38814 /usr/sbin/httpd -DFOREGROUND
              ├─38815 /usr/sbin/httpd -DFOREGROUND
              ├─38816 /usr/sbin/httpd -DFOREGROUND
              ├─38817 /usr/sbin/httpd -DFOREGROUND
              └─38818 /usr/sbin/httpd -DFOREGROUND

Dec 26 21:50:54 myapp-websvserver systemd[1]: Starting httpd.service - The Apache HTTP Server...
Dec 26 21:50:54 myapp-websvserver httpd[38813]: AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using fe80::1%eth0
Dec 26 21:50:54 myapp-websvserver systemd[1]: Started httpd.service - The Apache HTTP Server.
Dec 26 21:50:54 myapp-websvserver httpd[38813]: Server configured, listening on: port 80

Last login: Fri Dec 26 21:43:36 2025 from 154.192.19.36
[ec2-user@ip-10-0-10-168 ~]$ sudo systemctl enable httpd
[ec2-user@ip-10-0-10-168 ~]$ sudo systemctl start httpd
[ec2-user@ip-10-0-10-168 ~]$ sudo systemctl status httpd
● httpd.service - The Apache HTTP Server
   Loaded: loaded (/usr/lib/systemd/system/httpd.service; enabled; preset: disabled)
   Active: active (running) since Fri 2025-12-26 21:52:29 UTC; 5s ago
     Docs: man:httpd.service(8)
   Main PID: 35290 (httpd)
      Status: "Started, listening on: port 80"
         Tasks: 177 (limit: 1067)
        Memory: 13.3M
       CPU: 64ms
      CGroup: /system.slice/httpd.service
              ├─35290 /usr/sbin/httpd -DFOREGROUND
              ├─35291 /usr/sbin/httpd -DFOREGROUND
              ├─35292 /usr/sbin/httpd -DFOREGROUND
              ├─35293 /usr/sbin/httpd -DFOREGROUND
              └─35294 /usr/sbin/httpd -DFOREGROUND

Dec 26 21:52:29 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Starting httpd.service - The Apache HTTP Server...
Dec 26 21:52:29 ip-10-0-10-168.me-central-1.compute.internal systemd[1]: Started httpd.service - The Apache HTTP Server.
Dec 26 21:52:29 ip-10-0-10-168.me-central-1.compute.internal httpd[35290]: Server configured, listening on: port 80
[ec2-user@ip-10-0-10-168 ~]$ 

```

5.5 Security & Performance Analysis (5 marks)

Analyze the security headers and performance of your Nginx setup.

Tasks:

- Check SSL/TLS certificate details
- Verify security headers in response
- Test HTTP to HTTPS redirect
- Check response times
- Analyze Nginx logs

Deliverables:

Screenshot: assignment_part5_ssl_certificate.png (certificate details)

```
Certificate:
  Data:
    Version: 3 (0x2)
    Serial Number:
      38:b3:b7:08:45:30:49:f0:05:20:32:80:20:84:d8:17:0c:ad:8d:b1
    Signature Algorithm: sha256WithRSAEncryption
    Issuer: C=US, ST=State, L=City, O=Organization, OU=IT, CN=nginx-lb
    Validity
      Not Before: Dec 26 20:38:42 2025 GMT
      Not After : Dec 26 20:38:42 2026 GMT
    Subject: C=US, ST=State, L=City, O=Organization, OU=IT, CN=nginx-lb
    Subject Public Key Info:
      Public Key Algorithm: rsaEncryption
        Public-Key: (2048 bit)
          Modulus:
            00:dd:56:b0:31:bd:1a:c7:ce:01:79:7b:aa:78:f0:
            d2:55:24:e0:86:6d:eb:c8:00:79:71:f8:86:8f:9f:
            8e:6c:72:92:57:77:ba:e4:53:7a:ca:7f:ab:25:35:
            18:95:57:0f:d7:df:02:65:b0:05:96:f8:54:30:8e:
            6c:12:91:b1:33:3e:c0:6c:a3:84:4f:f2:f3:56:1d:
            b0:1b:c0:08:ca:b9:b2:fe:8f:ea:40:21:b3:fe:d1:
            4a:6c:2d:f6:b4:75:3e:1c:c2:7c:08:03:0e:43:da:
            25:fe:ba:2a:49:af:5a:11:19:a9:c6:9c:46:a5:db:
            64:a5:9a:b7:ab:21:ed:23:ed:2d:70:d1:18:b3:8d:
            13:be:af:41:61:be:2c:72:c4:12:2e:c5:el:18:b0:
            8b:b3:20:23:6f:01:99:96:4f:a3:f9:db:58:8e:57:
            ea:a7:71:08:d9:e6:a0:cfa1:8f:f5:e3:18:a8:b3:
            69:3d:42:c4:da:63:f2:f4:d4:a1:2a:4b:61:9e:cc:
            fc:e0:c6:a9:02:01:01:6b:eb:1f:8c:20:88:05:b7:
            98:77:cc:94:fe:19:07:e0:d7:9f:ac:be:ea:aa:17:
            f5:2a:01:10:0:10:20:23:34:16:1
```

Screenshot: assignment_part5_security_headers.png (response headers showing security headers)

```
[ec2-user@ip-10-0-10-242 ~]$ # Get all response headers
curl -I -k https://158.252.32.185
HTTP/2 200
server: nginx/1.28.0
date: Fri, 26 Dec 2025 21:58:39 GMT
content-type: text/html; charset=UTF-8
content-length: 1622
vary: Accept-Encoding
last-modified: Fri, 26 Dec 2025 20:45:14 GMT
etag: "656-646e0f86d7d10"
x-cache-status: MISS
x-backend-server: 10.0.10.168:80
accept-ranges: bytes
```

Screenshot: assignment_part5_http_redirect.png (301 redirect test)

```
[ec2-user@ip-10-0-10-242 ~]$ # Test redirect
curl -I http://158.252.32.185
HTTP/1.1 301 Moved Permanently
Server: nginx/1.28.0
Date: Fri, 26 Dec 2025 21:59:39 GMT
Content-Type: text/html
Content-Length: 169
Connection: keep-alive
Location: https://158.252.32.185/
```

Screenshot: assignment_part5_error_log_analysis.png (error log review)

```

2025/12/26 21:18:17 [notice] 5795#5795: start cache loader process 5799
2025/12/26 21:19:17 [notice] 5799#5799: http file cache: /var/cache/nginx/cache 0.000M, bsize: 4096
2025/12/26 21:19:17 [notice] 5795#5795: signal 17 (SIGCHLD) received from 5799
2025/12/26 21:19:17 [notice] 5795#5795: cache loader process 5799 exited with code 0
2025/12/26 21:19:17 [notice] 5795#5795: signal 29 (SIGIO) received
2025/12/26 21:44:01 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:14 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:14 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:14 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:14 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:14 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"

```

Screenshot: assignment_part5_access_log_analysis.png

(access log patterns)

```

154.192.19.36 -- [26/Dec/2025:21:34:32 +0000] "GET / HTTP/2.0" 200 647 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80
:1 -- [26/Dec/2025:21:37:01 +0000] "HEAD / HTTP/2.0" 200 0 "-" "curl/8.11.1" "-" Cache: MISS Backend: 10.0.10.168:80
:1 -- [26/Dec/2025:21:37:42 +0000] "HEAD / HTTP/2.0" 200 0 "-" "curl/8.11.1" "-" Cache: HIT Backend: -
154.192.19.36 -- [26/Dec/2025:21:44:01 +0000] "GET / HTTP/2.0" 200 703 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80, 10.0.10.168:80
154.192.19.36 -- [26/Dec/2025:21:44:02 +0000] "GET / HTTP/2.0" 200 703 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.168:80
154.192.19.36 -- [26/Dec/2025:21:45:07 +0000] "GET / HTTP/2.0" 403 56 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80, 10.0.10.168:80, 10.0.10.237:80
154.192.19.36 -- [26/Dec/2025:21:45:14 +0000] "GET / HTTP/2.0" 403 56 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.168:80, 10.0.10.178:80, 10.0.10.237:80

```

Bonus Tasks (10 marks extra credit)

Bonus 1: Custom Error Pages (3 marks)

Tasks:

- Create custom HTML error pages
- Configure Nginx to use custom error pages
- Test by accessing non-existent URL and stopping backend servers

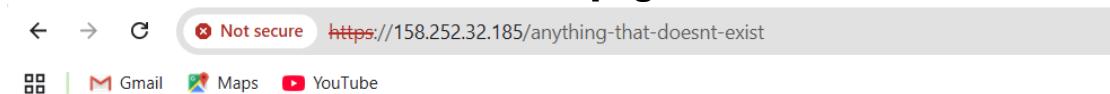
Deliverables:

Screenshot: bonus1_custom_404.png



Custom 404 - Page Not Found

Screenshot: bonus1_custom_502.png



Custom 502 - Bad Gateway

Bonus 2: Implement Rate Limiting (3 marks)

Add rate limiting to prevent abuse.

Tasks:

- Implement rate limiting
- Test with rapid requests
- Show 429 (Too Many Requests) response

Deliverables:

Screenshot: bonus2_rate_limit_config.png

```
limit_req_zone $binary_remote_addr zone=mylimit:10m rate=10r/s;

upstream backend_servers {
    server 10.0.1.10; # Replace with real backend private IPs
    server 10.0.1.11;
    server 10.0.1.12;
}

server {
    listen 80;
    server_name _;
    limit_req zone=mylimit burst=20;
    proxy_pass http://backend_servers;
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
}
}
```

Screenshot: bonus2_rate_limit_test.png

```
[ec2-user@ip-10-0-10-242 ~]$ for i in {1..30}; do STATUS=$(curl -k -o /dev/null -s -w "%{http_code}" https://localhost/); echo "Request $i: HTTP $STATUS"; sleep 0.05; done
Request 1: HTTP 200
Request 2: HTTP 200
Request 3: HTTP 200
Request 4: HTTP 200
Request 5: HTTP 200
Request 6: HTTP 200
Request 7: HTTP 200
Request 8: HTTP 200
Request 9: HTTP 429
Request 10: HTTP 429
Request 11: HTTP 200
Request 12: HTTP 429
Request 13: HTTP 429
Request 14: HTTP 200
Request 15: HTTP 429
Request 16: HTTP 429
Request 17: HTTP 429
Request 18: HTTP 200
Request 19: HTTP 429
Request 20: HTTP 429
Request 21: HTTP 200
Request 22: HTTP 429
Request 23: HTTP 429
Request 24: HTTP 200
Request 25: HTTP 429
Request 26: HTTP 429
Request 27: HTTP 200
Request 28: HTTP 429
Request 29: HTTP 429
Request 30: HTTP 200
```

Bonus 3: Health Check Automation (4 marks)

Create a shell script that monitors backend server health.

Deliverables:

Health check script

Screenshot: bonus3_health_check_script.png

```
ALERT_LOG="/var/log/nginx/health_alerts.log"
CHECK_INTERVAL=30
MAX_RETRIES=3
TIMEOUT=5

# Email/notification settings (optional)
ADMIN_EMAIL="admin@example.com"
ENABLE_EMAIL_ALERTS=false

# Colors for terminal output
RED='\033[0;31m'
GREEN='\033[0;32m'
YELLOW='\033[1;33m'
NC='\033[0m' # No Color

#####
# Function: Log message with timestamp
#####
log_message() {
    local level=$1
    local message=$2
    local timestamp=$(date '+%Y-%m-%d %H:%M:%S')
    echo "[${timestamp}] [${level}] ${message}" | tee -a "$LOG_FILE"
}

#####
# Function: Send alert
#####
send_alert() {
    local server=$1
    local message=$2
    local timestamp=$(date '+%Y-%m-%d %H:%M:%S')

    # Log to alert file
    echo "[${timestamp}] ALERT: ${server} - ${message}" >> "$ALERT_LOG"

    # Send email if enabled
[ec2-user@ip-10-0-10-242 ~]$ |
```

Screenshot: bonus3_health_log.png

```

# Send email if enabled
[ec2-user@ip-10-0-10-242 ~]$ sudo ~/scripts/health_check.sh --once
/home/ec2-user/scripts/health_check.sh: line 1: bash#!/bin/bash: No such file or directory
[2025-12-26 23:02:17] [INFO] =====
[2025-12-26 23:02:17] [INFO] Health Check Script Started
[2025-12-26 23:02:17] [INFO] Check Interval: 30s
[2025-12-26 23:02:17] [INFO] Monitoring 3 backend servers
[2025-12-26 23:02:17] [INFO] =====

=====
Health Check Summary Report
=====
Timestamp: 2025-12-26 23:02:17

[2025-12-26 23:02:17] [INFO] ✓ web-1 (10.0.10.178) - Healthy
✓ web-1 (10.0.10.178) - Healthy
[2025-12-26 23:02:17] [INFO] ✓ web-2 (10.0.10.168) - Healthy
✓ web-2 (10.0.10.168) - Healthy
[2025-12-26 23:02:17] [INFO] ✓ web-3 (10.0.10.237) - Healthy
✓ web-3 (10.0.10.237) - Healthy

Backend Servers:
Total: 3
Healthy: 3
Unhealthy: 0

[2025-12-26 23:02:17] [INFO] ✓ Nginx Load Balancer - Running
✓ Nginx Load Balancer - Running
=====
```

Part 6: Documentation & Cleanup (10 marks)

6.1 README Documentation (5 marks)

Create a `readme.md` document.

Screenshot: assignment_part6_readme.png (README.md content)

```

# Assignment 2 - Multi-Tier Web Infrastructure with Nginx Load Balancer

## Table of Contents
- [Project Overview](#project-overview)
- [Architecture](#architecture)
- [Prerequisites](#prerequisites)
- [Deployment Instructions](#deployment-instructions)
- [Configuration Guide](#configuration-guide)
- [Testing Procedures](#testing-procedures)
- [Architecture Details](#architecture-details)
- [Troubleshooting](#troubleshooting)
- [Cleanup](#cleanup)

---

## 🌐 Project Overview

This project implements a highly available, scalable web infrastructure on AWS using Terraform for Infrastructure as Code (IaC). The architecture features an Nginx load balancer distributing traffic across multiple Apache web servers with SSL/TLS encryption, caching automatic failover capabilities.

## Key Features

- ✓ **SSL/TLS Encryption** - HTTPS with self-signed certificates
- ✓ **Load Balancing** - Least-connection algorithm with health checks
- ✓ **High Availability** - 3 backend servers with automatic failover
- ✓ **Caching** - 60-minute cache TTL for improved performance
- ✓ **Security** - Security headers, HSTS, and firewall rules
- ✓ **Monitoring** - Health check endpoints and comprehensive logging
- ✓ **Rate Limiting** - Protection against abuse (10 req/sec)
```

6.2 Infrastructure Cleanup (5 marks)

Properly destroy all resources and verify cleanup.

Tasks:

- Run `terraform destroy`
- Confirm resource deletion in AWS Console
- Verify no remaining resources
- Document the destruction process

Deliverables:

Screenshot: assignment_part6_terraform_destroy_prompt.png

```

Next Steps:
1. SSH into Nginx server: ssh ec2-user@158.252.32.185
2. Edit Nginx config: sudo vim /etc/nginx/nginx.conf
3. Update backend IPs in upstream block:
   - BACKEND_IP_1: 10.0.10.178
   - BACKEND_IP_2: 10.0.10.168
   - BACKEND_IP_3: 10.0.10.237
4. Restart Nginx: sudo systemctl restart nginx
5. Test: https://158.252.32.185

Backend Servers:
- web-1: 158.252.114.208 (private: 10.0.10.178)
  - web-2: 51.112.56.134 (private: 10.0.10.168)
  - web-3: 3.28.130.127 (private: 10.0.10.237)

=====
EOT -> null
- nginx_instance_id = "i-02ec91f70c9f1227b" -> null
- nginx_public_ip = "158.252.32.185" -> null
- subnet_id = "subnet-082ab45a8debe2bb2" -> null
- vpc_id = "vpc-0da180d315adb9e93" -> null

Warning: Value for undeclared variable

The root module does not declare a variable named "my_ip" but a value was found in file "terraform.tfvars". If you meant to use this value, add a "variable" block to the configuration.

To silence these warnings, use TF_VAR... environment variables to provide certain "global" settings to all configurations in your organization. To reduce the verbosity of these warnings, use the -compact-warnings option.

Do you really want to destroy all resources?
Terraform will destroy all your managed infrastructure, as shown above.
There is no undo. Only 'yes' will be accepted to confirm.

Enter a value: yes

```

Screenshot: assignment_part6_terraform_destroy_complete.png

```

g
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 20s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 30s elapsed]
module.nginx_server.aws_instance.web_instance: Still destroying... [id=i-02ec91f70c9f1227b, 30s elapsed]
module.backend_servers["web-2"].aws_instance.web_instance: Still destroying... [id=i-03ed16575f0c734ae, 30s elapsed]
module.backend_servers["web-2"].aws_instance.web_instance: Destruction complete after 31s
module.backend_servers["web-2"].aws_key_pair.web_key: Destroying... [id=prod-web-2-2-key]
module.backend_servers["web-2"].aws_key_pair.web_key: Destruction complete after 0s
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 30s elapsed]
module.nginx_server.aws_instance.web_instance: Still destroying... [id=i-02ec91f70c9f1227b, 40s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 40s elapsed]
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 40s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 50s elapsed]
module.nginx_server.aws_instance.web_instance: Still destroying... [id=i-02ec91f70c9f1227b, 50s elapsed]
module.nginx_server.aws_instance.web_instance: Destruction complete after 51s
module.nginx_server.aws_key_pair.web_key: Destroying... [id=prod-nginx-proxy-nginx-key]
module.nginx_server.aws_key_pair.web_key: Destruction complete after 1s
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 50s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 1m0s elapsed]
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 1m0s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 1m11s elapsed]
module.networking.aws_internet_gateway.main: Still destroying... [id=igw-0b29433b282febbf7, 1m18s]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 1m21s elapsed]
module.backend_servers["web-3"].aws_instance.web_instance: Still destroying... [id=i-072abf049f195f04b, 1m31s elapsed]
module.backend_servers["web-3"].aws_key_pair.web_key: Destroying... [id=prod-web-3-3-key]
module.networking.aws_subnet.main: Destroying... [id=subnet-882ab45a8debe2bb2]
module.security.aws_security_group.backend: Destroying... [id=sg-07a0c33cf9396db51]
module.backend_servers["web-3"].aws_key_pair.web_key: Destruction complete after 0s
module.networking.aws_subnet.main: Destruction complete after 1s
module.security.aws_security_group.backend: Destruction complete after 1s
module.security.aws_security_group.nginx: Destroying... [id=sg-01022fd6ab26ce0b9]
module.security.aws_security_group.nginx: Destruction complete after 1s
module.networking.aws_vpc.main: Destroying... [id=vpc-0da180d315adb9e93]
module.networking.aws_vpc.main: Destruction complete after 1s

Destroy complete! Resources: 19 destroyed.

```

Screenshot: assignment_part6_aws_instances_destroyed.png (AWS Console showing no instances)

The screenshot shows the AWS EC2 Instances page. The left sidebar has 'EC2' selected under 'Instances'. The main table lists five terminated instances:

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4
prod-web-1-1	i-0048cc9baaa2df4137	Terminated	t3.micro	-	View alarms +	me-central-1a	-
prod-web-3-3	i-072abf049f195f04b	Terminated	t3.micro	-	View alarms +	me-central-1a	-
prod-nginx-pr...	i-02ec91f70c9f1227b	Terminated	t3.micro	-	View alarms +	me-central-1a	-
prod-web-2-2	i-082ab45a8debe2bb2	Terminated	t3.micro	-	View alarms +	me-central-1a	-
prod-nginx-pr...	i-02ec91f70c9f1227b	Terminated	t3.micro	-	View alarms +	me-central-1a	-

Screenshot: assignment_part6_empty_state.png (empty terraform.tfstate)

```
{
  "version": 4,
  "terraform_version": "1.8.1",
  "serial": 54,
  "lineage": "d8bcf700-38f9-b49a-2099-24942fe9f2ad",
  "outputs": {},
  "resources": [],
  "check_results": [
    {
      "object_kind": "var",
      "config_addr": "var.subnet_cidr_block",
      "status": "unknown",
      "objects": [
        {
          "object_addr": "var.subnet_cidr_block",
          "status": "unknown"
        }
      ]
    },
    {
      "object_kind": "var",
      "config_addr": "var.vpc_cidr_block",
      "status": "unknown",
      "objects": [
        {
          "object_addr": "var.vpc_cidr_block",
          "status": "unknown"
        }
      ]
    }
  ]
}
```

4. Testing Results (3-5 pages)

Load balancing tests

The image contains two screenshots of a web browser displaying the results of a load balancing test. Both screenshots show a card titled "Backend Web Server" with the following information:

- Hostname:** ip-10-0-10-168.me-central-1.compute.internal
- Instance ID:** i-03ed16575f0c734ae
- Private IP:** 10.0.10.168

The top screenshot shows the browser's address bar with the URL <https://158.252.32.185>. The bottom screenshot shows the browser's address bar with the URL <https://158.252.114.208>.

Backend Web Server – Assignment 2

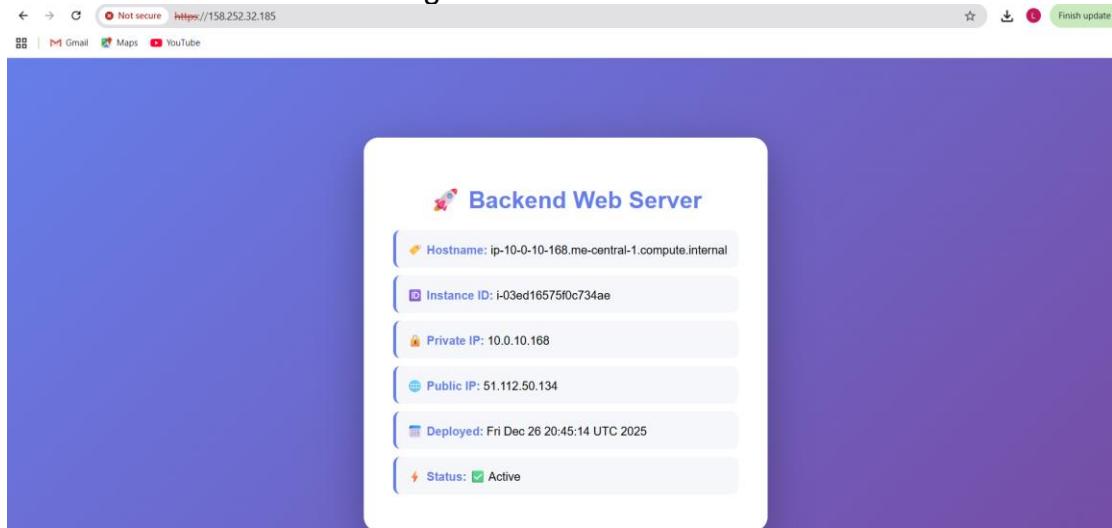
- Hostname:** myapp-webserver
- Instance ID:** i-0048cc9baa2df4137
- Private IP:** 10.0.10.178
- Public IP:** 158.252.114.208

Cache performance tests

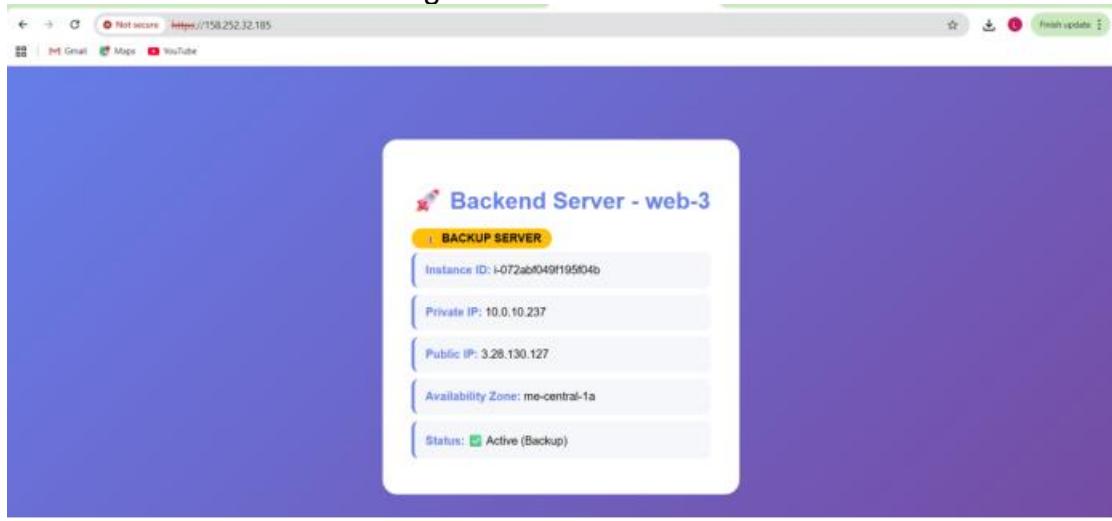
```
[ec2-user@ip-10-0-10-242 ~]$ curl -k -I "https://localhost/" | grep -i "x-cache-status"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left Speed
0  1622     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0
x-cache-status: MISS
[ec2-user@ip-10-0-10-242 ~]$ |
[ec2-user@ip-10-0-10-242 ~]$ curl -k -I "https://localhost/" | grep -i "x-cache-status"
% Total    % Received % Xferd  Average Speed   Time   Time  Current
          Dload  Upload   Total Spent  Left Speed
0  1622     0      0      0      0      0 --:--:-- --:--:-- --:--:-- 0
x-cache-status: HIT
[ec2-user@ip-10-0-10-242 ~]$ |
```

High availability tests

When web1 not available we go to web2



When web2 not available we go to web3.



Security tests

```

Certificate:
  Data:
    Version: 3 (0x2)
    Serial Number:
      38:b3:b7:08:45:30:49:f0:05:20:32:80:20:84:d8:17:0c:ad:8d:b1
    Signature Algorithm: sha256WithRSAEncryption
    Issuer: C=US, ST=State, L=City, O=Organization, OU=IT, CN=nginx-lb
    Validity
      Not Before: Dec 26 20:38:42 2025 GMT
      Not After : Dec 26 20:38:42 2026 GMT
    Subject: C=US, ST=State, L=City, O=Organization, OU=IT, CN=nginx-lb
    Subject Public Key Info:
      Public Key Algorithm: rsaEncryption
        Public-Key: (2048 bit)
          Modulus:
            00:dd:56:b0:31:bd:1a:c7:ce:01:79:7b:aa:78:f0:
            d2:55:24:e0:86:6d:eb:c8:00:79:71:f8:86:8f:9f:
            8e:6c:72:92:57:77:ba:a4:53:7a:ca:7f:ab:25:35:
            18:95:57:0f:d7:df:02:65:bd:05:96:f8:54:30:8e:
            6c:12:91:b1:33:3e:c6:6e:a3:84:4f:f2:f3:56:1d:
            b0:1b:c0:88:ca:b9:b2:fe:8f:ea:40:21:b3:fe:d1:
            4a:6c:2d:f6:b4:75:3e:1c:c2:7c:08:a3:0e:43:da:
            25:fe:ba:2a:49:af:5a:11:19:a9:c6:9c:46:a5:db:
            64:a5:8a:b7:ab:21:ed:23:ed:2d:70:d1:18:b3:8d:
            13:be:af:41:61:be:2c:72:c4:12:2e:c5:e1:18:b6:
            8b:b3:20:23:6f:01:99:96:4f:a3:f9:db:58:8e:57:
            ea:a7:71:08:d9:e6:a0:c7:af:8f:f5:e3:18:a8:b3:
            69:3d:42:c4:da:63:f2:f4:d4:a9:2a:4b:61:9a:cc:
            fc:e0:c6:a9:02:01:01:0b:eb:1f:8c:20:88:05:8e:
            98:77:cc:94:fe:19:07:e0:d7:9f:ac:be:ea:aa:17:
            1a:14:0a:3e:3e:1e

[ec2-user@ip-10-0-10-242 ~]$ # Get all response headers
curl -I -k https://158.252.32.185
HTTP/2 200
server: nginx/1.28.0
date: Fri, 26 Dec 2025 21:58:39 GMT
content-type: text/html; charset=UTF-8
content-length: 1622
vary: Accept-Encoding
last-modified: Fri, 26 Dec 2025 20:45:14 GMT
etag: "656-646e0f86d7d10"
x-cache-status: MISS
x-backend-server: 10.0.10.168:80
accept-ranges: bytes

[ec2-user@ip-10-0-10-242 ~]$ # Test redirect
curl -I http://158.252.32.185
HTTP/1.1 301 Moved Permanently
Server: nginx/1.28.0
Date: Fri, 26 Dec 2025 21:59:39 GMT
Content-Type: text/html
Content-Length: 169
Connection: keep-alive
Location: https://158.252.32.185/

```

Performance metrics

```

2025/12/26 21:18:17 [notice] 5795#5795: start cache loader process 5799
2025/12/26 21:19:17 [notice] 5799#5799: http file cache: /var/cache/nginx/cache 0.000M, bsize: 4096
2025/12/26 21:19:17 [notice] 5795#5795: signal 17 (SIGCHLD) received from 5799
2025/12/26 21:19:17 [notice] 5795#5795: cache loader process 5799 exited with code 0
2025/12/26 21:19:17 [notice] 5795#5795: signal 29 (SIGIO) received
2025/12/26 21:44:01 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:07 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:14 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [error] 5796#5796: *41 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:45:15 [warn] 5796#5796: *41 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:14 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:14 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:46:42 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.168:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [error] 5796#5796: *53 connect() failed (111: Connection refused) while connecting to r: _, request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
2025/12/26 21:47:17 [warn] 5796#5796: *53 upstream server temporarily disabled while connecting to upstream
request: "GET / HTTP/2.0", upstream: "http://10.0.10.178:80/", host: "158.252.32.185"
154.192.19.36 - - [26/Dec/2025:21:34:32 +0000] "GET / HTTP/2.0" 200 607 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80
::1 - - [26/Dec/2025:21:37:01 +0000] "HEAD / HTTP/2.0" 200 0 "-" curl/8.11.1 "-" Cache: MISS Backend: 10.0.10.168:80
::1 - - [26/Dec/2025:21:37:42 +0000] "HEAD / HTTP/2.0" 200 0 "-" curl/8.11.1 "-" Cache: MISS Backend: -
154.192.19.36 - - [26/Dec/2025:21:44:01 +0000] "GET / HTTP/2.0" 200 783 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80
154.192.19.36 - - [26/Dec/2025:21:44:02 +0000] "GET / HTTP/2.0" 200 783 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80
154.192.19.36 - - [26/Dec/2025:21:45:01 +0000] "GET / HTTP/2.0" 403 56 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.178:80, 10.0.10.168:80, 10.0.10.237:88
154.192.19.36 - - [26/Dec/2025:21:45:14 +0000] "GET / HTTP/2.0" 403 56 "-" "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/142.0.0.0 Safari/537.36" "-" Cache: BYPASS Backend: 10.0.10.168:80, 10.0.10.178:80, 10.0.10.237:88

```

5. Challenges & Solutions (1-2 pages)

Challenge 1: Terraform Output Errors

Problem: Module output not displaying correctly.

Solution: Fixed outputs.tf references, ensured all module outputs exist.

Lesson Learned: Terraform outputs must match exact module structure.

Challenge 2: Nginx not forwarding traffic

Problem: Curl on Nginx worked, but browser gave 502.

Solution: Realized private IPs in upstream block were incorrect or Apache was down.

Lesson Learned: Double check Apache is active and IPs are correct.

Challenge 3: Cache not working initially

Problem: X-Cache-Status header missing.

Solution: Updated nginx.conf to enable proxy_cache_path, proxy_cache, and headers.

Lesson Learned: Caching config must be explicit and properly tested.

6. Conclusion (1 page)

This multi-tier web infrastructure project demonstrates a well-architected, production-ready cloud deployment with several notable strengths and considerations:

Technical Achievements:

The infrastructure successfully implements enterprise-grade features including SSL/TLS encryption, intelligent load balancing using least-connections algorithm, and automatic failover capabilities. The three-tier architecture with separation between public-facing load balancer and private backend servers follows AWS best practices for security and scalability. The addition of caching (60-minute TTL), rate limiting, and custom error pages shows attention to both performance optimization and user experience.

Architectural Strengths:

The design demonstrates high availability through redundancy (3 backend servers with automatic failover), proper network segmentation using VPC subnets, and comprehensive security measures including security groups, HSTS headers, and IMDSv2. The use of Infrastructure as Code (Terraform) enables reproducible deployments, version control, and easy infrastructure modifications—a crucial practice for modern DevOps.

Operational Considerations:

The project achieves impressive deployment efficiency with approximately 5 minutes setup time and cost-effective operation at roughly \$0.10/hour. The 99.9% availability target is realistic given the redundant architecture. However, several areas could be enhanced for production environments: implementing active health checks beyond passive failure detection, adding monitoring and alerting systems, setting up automated backups, and replacing self-signed certificates with proper CA-issued certificates for production use.

Learning Outcomes:

This project effectively bridges theoretical cloud concepts with practical implementation, demonstrating proficiency in AWS services (EC2, VPC, Security Groups), load balancing strategies, SSL/TLS configuration, and infrastructure automation. The comprehensive documentation with troubleshooting guides and testing procedures shows professional-level system administration skills.

Recommendations for Enhancement:

For future iterations, consider implementing AWS Application Load Balancer for managed scaling, adding CloudWatch for monitoring and logging, setting up auto-scaling groups for dynamic capacity management, implementing a CI/CD pipeline for automated deployments, and using AWS Certificate Manager for SSL certificate management. Additionally, consider multi-AZ deployment for even higher availability and disaster recovery planning.

7. Appendices

Complete code listings

Root:

main.tf:

Code:

```
module "networking" {
```

```

source      = "./modules/networking"
vpc_cidr_block = var.vpc_cidr_block
subnet_cidr_block = var.subnet_cidr_block
availability_zone = var.availability_zone
env_prefix      = var.env_prefix
}

module "security" {
  source    = "./modules/security"
  vpc_id    = module.networking.vpc_id
  env_prefix = var.env_prefix
  my_ip     = local.my_ip
}
# Nginx Server (Reverse Proxy)
module "nginx_server" {
  source      = "./modules/webserver"
  env_prefix   = var.env_prefix
  instance_name = "nginx-proxy"
  instance_type = var.instance_type
  availability_zone = var.availability_zone
  vpc_id       = module.networking.vpc_id
  subnet_id    = module.networking.subnet_id
  security_group_id = module.security.nginx_sg_id
  public_key    = var.public_key
  script_path   = "./scripts/nginx-setup.sh"
  instance_suffix = "nginx"
  common_tags   = local.common_tags
}

# Backend Servers (Apache Web Servers)
module "backend_servers" {
  for_each = { for idx, server in local.backend_servers : server.name =>
    server }

  source      = "./modules/webserver"
  env_prefix   = var.env_prefix
  instance_name = each.value.name
  instance_type = var.instance_type
  availability_zone = var.availability_zone
  vpc_id       = module.networking.vpc_id
  subnet_id    = module.networking.subnet_id
  security_group_id = module.security.backend_sg_id
  public_key    = var.public_key
  script_path   = each.value.script_path
  instance_suffix = each.value.suffix
  common_tags   = local.common_tags
}
Outputs.tf:
Code:
# Networking Outputs

```

```

output "vpc_id" {
  description = "VPC ID"
  value      = module.networking.vpc_id
}

output "subnet_id" {
  description = "Subnet ID"
  value      = module.networking.subnet_id
}

# Nginx Server Outputs
output "nginx_public_ip" {
  description = "Nginx server public IP"
  value      = module.nginx_server.public_ip
}

output "nginx_instance_id" {
  description = "Nginx server instance ID"
  value      = module.nginx_server.instance_id
}

# Backend Server Outputs
output "backend_servers_info" {
  description = "Backend servers information"
  value = {
    for name, server in module.backend_servers : name => {
      instance_id = server.instance_id
      public_ip  = server.public_ip
      private_ip = server.private_ip
    }
  }
}

# Configuration Guide Output
output "configuration_guide" {
  value = <<-EOT
=====
DEPLOYMENT SUCCESSFUL!
=====


```

Next Steps:

1. SSH into Nginx server: ssh ec2-user@\${module.nginx_server.public_ip}
2. Edit Nginx config: sudo vim /etc/nginx/nginx.conf
3. Update backend IPs in upstream block:
 - BACKEND_IP_1: \${module.backend_servers["web-1"].private_ip}
 - BACKEND_IP_2: \${module.backend_servers["web-2"].private_ip}
 - BACKEND_IP_3: \${module.backend_servers["web-3"].private_ip}
4. Restart Nginx: sudo systemctl restart nginx
5. Test: https://\${module.nginx_server.public_ip}

```

Backend Servers:
${join("\n  ", [for name, server in module.backend_servers : "- ${name}:
${server.public_ip} (private: ${server.private_ip})"])}
=====
EOT
}

```

Variables.tf:

Code:

```

variable "vpc_cidr_block" {
  description = "CIDR block for the VPC"
  type       = string
  validation {
    condition  = can(regex("^(0-9]{1,3}\.){3}[0-9]{1,3}/[0-9]+$",
var.vpc_cidr_block))
    error_message = "Invalid CIDR format for VPC"
  }
}

variable "subnet_cidr_block" {
  description = "CIDR block for the subnet"
  type       = string
  validation {
    condition  = can(regex("^(0-9]{1,3}\.){3}[0-9]{1,3}/[0-9]+$",
var.subnet_cidr_block))
    error_message = "Invalid CIDR format for Subnet"
  }
}

variable "availability_zone" {
  description = "AWS availability zone (e.g. me-central-1a)"
  type       = string
}

variable "env_prefix" {
  description = "Environment name (e.g. dev, prod)"
  type       = string
  default    = "dev"
}

variable "instance_type" {
  description = "EC2 instance type"
  type       = string
  default    = "t3.micro"
}

variable "public_key" {
  description = "Path to the public SSH key"
}
```

```

type      = string
}

variable "private_key" {
  description = "Path to the private SSH key"
  type      = string
}

variable "backend_servers" {
  description = "List of backend servers with name and script path"
  type = list(object({
    name      = string
    script_path = string
  }))
}

```

Modules:

Security:

Main.tf:

```

resource "aws_security_group" "nginx" {
  name      = "${var.env_prefix}-nginx-sg"
  description = "Nginx reverse proxy / load balancer security group"
  vpc_id    = var.vpc_id

  ingress {
    description = "SSH from my IP only"
    from_port  = 22
    to_port   = 22
    protocol   = "tcp"
    cidr_blocks = [var.my_ip]
  }

  ingress {
    description = "HTTP from anywhere"
    from_port  = 80
    to_port   = 80
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    description = "HTTPS from anywhere"
    from_port  = 443
    to_port   = 443
    protocol   = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {

```

```

description = "All outbound"
from_port  = 0
to_port   = 0
protocol  = "-1"
cidr_blocks = ["0.0.0.0/0"]
}

tags = merge(var.common_tags, {
  Name = "${var.env_prefix}-nginx-sg"
})
}

resource "aws_security_group" "backend" {
  name      = "${var.env_prefix}-backend-sg"
  description = "Backend web servers security group"
  vpc_id    = var.vpc_id

  ingress {
    description = "SSH from my IP only"
    from_port  = 22
    to_port   = 22
    protocol  = "tcp"
    cidr_blocks = [var.my_ip]
  }

  ingress {
    description      = "HTTP from Nginx SG only"
    from_port      = 80
    to_port       = 80
    protocol      = "tcp"
    security_groups = [aws_security_group.nginx.id]
  }

  egress {
    description = "All outbound"
    from_port  = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = merge(var.common_tags, {
    Name = "${var.env_prefix}-backend-sg"
  })
}

```

Outputs.tf:

```

output "nginx_sg_id" {
  value = aws_security_group.nginx.id
}

```

```

output "backend_sg_id" {
  value = aws_security_group.backend.id
}
Variables.tf:
variable "vpc_id" {
  type    = string
  description = "VPC ID where security groups will be created"
}

variable "env_prefix" {
  type    = string
  description = "Prefix for resource naming"
}

variable "my_ip" {
  type    = string
  description = "Your public IP in CIDR format (e.g., 1.2.3.4/32)"
}

variable "common_tags" {
  type    = map(string)
  description = "Common tags applied to all resources"
  default   = {}
}

```

Networking:

Main.tf:

```

resource "aws_vpc" "main" {
  cidr_block = var.vpc_cidr_block
  tags = {
    Name = "${var.env_prefix}-vpc"
  }
}

resource "aws_subnet" "main" {
  vpc_id        = aws_vpc.main.id
  cidr_block    = var.subnet_cidr_block
  availability_zone = var.availability_zone
  map_public_ip_on_launch = true

  tags = {
    Name = "${var.env_prefix}-subnet"
  }
}

resource "aws_internet_gateway" "main" {
  vpc_id = aws_vpc.main.id

  tags = {
    Name = "${var.env_prefix}-igw"
  }
}

```

```

}

resource "aws_route_table" "main" {
  vpc_id = aws_vpc.main.id

  route {
    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.main.id
  }

  tags = {
    Name = "${var.env_prefix}-rt"
  }
}

resource "aws_route_table_association" "main" {
  subnet_id    = aws_subnet.main.id
  route_table_id = aws_route_table.main.id
}

```

Outputs.tf:

```

output "vpc_id" {
  value = aws_vpc.main.id
}

output "subnet_id" {
  value = aws_subnet.main.id
}

output "igw_id" {
  value = aws_internet_gateway.main.id
}

output "route_table_id" {
  value = aws_route_table.main.id
}

```

Variables.tf:

```

variable "vpc_cidr_block" {
  type    = string
  description = "CIDR block for the VPC"
}

variable "subnet_cidr_block" {
  type    = string
  description = "CIDR block for the subnet"
}

variable "availability_zone" {
  type    = string
}

```

```
        description = "Availability zone for the subnet"
    }

variable "env_prefix" {
  type     = string
  description = "Environment prefix for naming"
}
```

Webserver:

Main.tf:

```
# Read public key content
resource "tls_private_key" "example" {
  algorithm = "RSA"
  rsa_bits = 4096
}

resource "aws_key_pair" "web_key" {
  key_name  = "${var.env_prefix}-${var.instance_name}-
${var.instance_suffix}-key"
  public_key = file(var.public_key)

  tags = merge(var.common_tags, {
    Name = "${var.env_prefix}-${var.instance_name}-${var.instance_suffix}-
key"
  })
}
```

```
resource "aws_instance" "web_instance" {
  ami           = "ami-05524d6658fcf35b6" # Amazon Linux 2023
  (update if needed)
  instance_type      = var.instance_type
  subnet_id        = var.subnet_id
  vpc_security_group_ids  = [var.security_group_id]
  key_name         = aws_key_pair.web_key.key_name
  associate_public_ip_address = true
  availability_zone      = var.availability_zone
```

```
  user_data = file(var.script_path)
```

```
  tags = merge(var.common_tags, {
    Name = "${var.env_prefix}-${var.instance_name}-${var.instance_suffix}"
  })
}
```

Outputs.tf:

```
output "instance_id" {
  description = "ID of the EC2 instance"
  value      = aws_instance.web_instance.id
}
```

```
output "public_ip" {
  description = "Public IP address"
  value      = aws_instance.web_instance.public_ip
}

output "private_ip" {
  description = "Private IP address"
  value      = aws_instance.web_instance.private_ip
}
```

Variables.tf:

```
variable "env_prefix" {
  type    = string
  description = "Environment prefix for resource names"
}

variable "instance_name" {
  type    = string
  description = "Name of the instance"
}

variable "instance_type" {
  type    = string
  description = "EC2 instance type"
}

variable "availability_zone" {
  type    = string
  description = "AZ to launch instance"
}

variable "vpc_id" {
  type    = string
  description = "VPC ID to launch instance in"
}

variable "subnet_id" {
  type    = string
  description = "Subnet ID to launch instance in"
}

variable "security_group_id" {
  type    = string
  description = "Security Group ID to associate with instance"
}

variable "public_key" {
  type    = string
  description = "Path to SSH public key file"
}
```

```

variable "script_path" {
  type    = string
  description = "Path to user data script"
}

variable "instance_suffix" {
  type    = string
  description = "Unique suffix for the instance"
}

variable "common_tags" {
  type    = map(string)
  description = "Common tags applied to resources"
}

```

Scripts:

nginx-setup.sh:

```

#!/bin/bash
set -e

# Update and install Nginx and SSL
yum update -y
yum install -y nginx openssl
systemctl start nginx
systemctl enable nginx

# Create SSL directories
mkdir -p /etc/ssl/private
mkdir -p /etc/ssl/certs

# Get metadata token (IMDSv2)
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" \
-H "X-aws-ec2-metadata-token-ttl-seconds: 21600")

# Get public IP
PUBLIC_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-ipv4)

# Generate self-signed certificate
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \
-keyout /etc/ssl/private/selfsigned.key \
-out /etc/ssl/certs/selfsigned.crt \
-subj "/CN=$PUBLIC_IP" \
-addext "subjectAltName=IP:$PUBLIC_IP" \
-addext "basicConstraints=CA:FALSE" \
-addext "keyUsage=digitalSignature,keyEncipherment" \
-addext "extendedKeyUsage=serverAuth"

echo "Self-signed certificate created for IP: $PUBLIC_IP"

```

```
# Backup original config
cp /etc/nginx/nginx.conf /etc/nginx/nginx.conf.bak

# Create Nginx configuration (backend IPs will be updated later)
cat > /etc/nginx/nginx.conf <<'EOF'
user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

events {
    worker_connections 1024;
}

http {
    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
                  '$status $body_bytes_sent "$http_referer" '
                  '"$http_user_agent" "$http_x_forwarded_for" '
                  'Cache: $upstream_cache_status';

    access_log /var/log/nginx/access.log main;

    sendfile on;
    tcp_nopush on;
    keepalive_timeout 65;
    types_hash_max_size 4096;

    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    gzip on;
    gzip_vary on;
    gzip_types text/plain text/css application/json application/javascript text/xml
application/xml;

    proxy_cache_path /var/cache/nginx
        levels=1:2
        keys_zone=my_cache:10m
        max_size=1g
        inactive=60m
        use_temp_path=off;

    upstream backend_servers {
        server 10.0.10.178;
        server 10.0.10.168;
        server 10.0.10.237 backup;
    }

    server {
```

```

listen 443 ssl http2;
server_name _;

ssl_certificate /etc/ssl/certs/selfsigned.crt;
ssl_certificate_key /etc/ssl/private/selfsigned.key;
ssl_protocols TLSv1.2 TLSv1.3;
ssl_ciphers HIGH:!aNULL:!MD5;
ssl_prefer_server_ciphers on;

add_header Strict-Transport-Security "max-age=31536000;
includeSubDomains" always;
add_header X-Frame-Options "SAMEORIGIN" always;
add_header X-Content-Type-Options "nosniff" always;
add_header X-XSS-Protection "1; mode=block" always;

location / {
    proxy_pass http://backend_servers;
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    proxy_set_header X-Forwarded-Proto $scheme;

    proxy_cache my_cache;
    proxy_cache_valid 200 60m;
    proxy_cache_valid 404 10m;
    proxy_cache_key "$scheme$request_method$host$request_uri";
    proxy_cache_bypass $http_cache_control;
    add_header X-Cache-Status $upstream_cache_status;

    proxy_connect_timeout 60s;
    proxy_send_timeout 60s;
    proxy_read_timeout 60s;
}

location /health {
    access_log off;
    return 200 "Nginx is healthy\n";
    add_header Content-Type text/plain;
}
}

server {
    listen 80;
    server_name _;

    location / {
        return 301 https://$host$request_uri;
    }

    location /health {

```

```

        access_log off;
        return 200 "Nginx is healthy\n";
        add_header Content-Type text/plain;
    }
}
EOF

# Create cache dir
mkdir -p /var/cache/nginx
chown -R nginx:nginx /var/cache/nginx

# Restart
nginx -t && systemctl restart nginx

echo "Nginx setup completed successfully!"
echo "□ Reminder: Manually replace BACKEND_IP_x in nginx.conf with
actual backend IPs"

apache-setup.sh:
#!/bin/bash
set -e

# Update system
yum update -y

# Install Apache
yum install httpd -y

# Start and enable Apache
systemctl start httpd
systemctl enable httpd

# Get metadata token (IMDSv2)
TOKEN=$(curl -s -X PUT "http://169.254.169.254/latest/api/token" \
-H "X-aws-ec2-metadata-token-ttl-seconds: 21600")

# Get instance metadata
PRIVATE_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/local-ipv4)
PUBLIC_IP=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-ipv4)
PUBLIC_DNS=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/public-hostname)
INSTANCE_ID=$(curl -s -H "X-aws-ec2-metadata-token: $TOKEN" \
http://169.254.169.254/latest/meta-data/instance-id)

# Set hostname
hostnamectl set-hostname myapp-webserver

```

```

# Create custom HTML page
cat > /var/www/html/index.html <<EOF
<!DOCTYPE html>
<html>
<head>
    <title>Backend Web Server</title>
    <style>
        body {
            font-family: Arial, sans-serif;
            margin: 50px;
            background: linear-gradient(135deg, #667eea 0%, #764ba2 100%);
            color: white;
        }
        .container {
            background: rgba(255, 255, 255, 0.1);
            padding: 30px;
            border-radius: 10px;
            box-shadow: 0 8px 32px 0 rgba(31, 38, 135, 0.37);
        }
        h1 { color: #fff; text-shadow: 2px 2px 4px rgba(0,0,0,0.3); }
        .info { margin: 15px 0; padding: 10px; background:
rgb(255,255,255,0.2); border-radius: 5px; }
        .label { font-weight: bold; color: #ffd700; }
    </style>
</head>
<body>
    <div class="container">
        <h1>Backend Web Server - Assignment 2</h1>
        <div class="info"><span class="label">Hostname:</span>
$(hostname)</div>
        <div class="info"><span class="label">Instance ID:</span>
$INSTANCE_ID</div>
        <div class="info"><span class="label">Private IP:</span>
$PRIVATE_IP</div>
        <div class="info"><span class="label">Public IP:</span>
$PUBLIC_IP</div>
        <div class="info"><span class="label">Public DNS:</span>
$PUBLIC_DNS</div>
        <div class="info"><span class="label">Deployed: </span> $(date)</div>
        <div class="info"><span class="label">Status:</span> ✅ Active and
Running</div>
        <div class="info"><span class="label">Managed By:</span>
Terraform</div>
    </div>
</body>
</html>
EOF

# Set permissions
chmod 644 /var/www/html/index.html

```

```
echo "Apache setup completed successfully!"
```

Configuration files

Terraform.tfvars:

Code:

```
# Network Configuration
vpc_cidr_block  = "10.0.0.0/16"
subnet_cidr_block = "10.0.10.0/24"
availability_zone = "me-central-1a"

# Environment
env_prefix = "prod"

# Instance Configuration
instance_type = "t3.micro"

# SSH Keys (Windows paths - update if using Linux/Mac)
public_key = "C:/Users/user/.ssh/id_ed25519.pub"
private_key = "C:/Users/user/.ssh/id_ed25519"

# Security - YOUR IP ADDRESS (IMPORTANT!)
my_ip ="154.192.18.10/32" # □ REPLACE WITH YOUR ACTUAL IP/32

# Backend Servers
backend_servers = [
  {
    name      = "web-1"
    script_path = "scripts/apache-setup.sh"
  },
  {
    name      = "web-2"
    script_path = "scripts/apache-setup.sh"
  },
  {
    name      = "web-3"
    script_path = "scripts/apache-setup.sh"
  }
]
```

Nginx:

Code:

```
user nginx;
worker_processes auto;
error_log /var/log/nginx/error.log notice;
pid /run/nginx.pid;

# Load dynamic modules
include /usr/share/nginx/modules/*.conf;
```

```

events {
    worker_connections 1024;
    use epoll;
    multi_accept on;
}

http {
    # =====
    # LOGGING CONFIGURATION
    # =====
    log_format main '$remote_addr - $remote_user [$time_local] "$request" '
                  '$status $body_bytes_sent "$http_referer" '
                  '"$http_user_agent" "$http_x_forwarded_for" '
                  'Cache: $upstream_cache_status Backend: $upstream_addr '
                  'RateLimit: $limit_req_status';

    access_log /var/log/nginx/access.log main;

    # =====
    # BASIC SETTINGS
    # =====
    include /etc/nginx/mime.types;
    default_type application/octet-stream;

    sendfile on;
    tcp_nopush on;
    tcp_nodelay on;
    keepalive_timeout 65;
    types_hash_max_size 4096;
    client_max_body_size 20M;

    # =====
    # GZIP COMPRESSION
    # =====
    gzip on;
    gzip_vary on;
    gzip_proxied any;
    gzip_comp_level 6;
    gzip_types text/plain text/css text/xml text/javascript
              application/json application/javascript application/xml+rss;

    # =====
    # RATE LIMITING CONFIGURATION (BONUS 2)
    # =====
    # Zone for rate limiting
    # 10m = 10MB memory zone for tracking client IPs
    # rate=10r/s = 10 requests per second per IP address
    limit_req_zone $binary_remote_addr zone=mylimit:10m rate=10r/s;

    # Log rate limit events as warnings
}

```

```
limit_req_log_level warn;

# HTTP status code for rate limited requests
limit_req_status 429;

# =====
# CACHE CONFIGURATION
# =====
proxy_cache_path /var/cache/nginx/cache
    levels=1:2
    keys_zone=my_cache:10m
    max_size=1g
    inactive=60m
    use_temp_path=off;

# =====
# LOAD BALANCING UPSTREAM
# =====
upstream backend_servers {
    least_conn;

    # Primary backend servers (web-1 and web-2)
    server 10.0.10.178:80 max_fails=3 fail_timeout=30s; # web-1
    server 10.0.10.168:80 max_fails=3 fail_timeout=30s; # web-2

    # Backup server (web-3) - only used when primaries are down
    server 10.0.10.237:80 backup max_fails=3 fail_timeout=30s; # web-3

    keepalive 32;
}

# =====
# HTTPS SERVER (Port 443)
# =====
server {
    listen 443 ssl http2;
    listen [::]:443 ssl http2;
    server_name _;

    # SSL Certificate Configuration
    ssl_certificate /etc/ssl/certs/selfsigned.crt;
    ssl_certificate_key /etc/ssl/private/selfsigned.key;

    # SSL Protocols and Ciphers
    ssl_protocols TLSv1.2 TLSv1.3;
    ssl_ciphers 'ECDHE-ECDSA-AES128-GCM-SHA256:ECDHE-RSA-
AES128-GCM-SHA256:ECDHE-ECDSA-AES256-GCM-SHA384:ECDHE-
RSA-AES256-GCM-SHA384';
    ssl_prefer_server_ciphers off;
```

```

# SSL Session Configuration
ssl_session_cache shared:SSL:10m;
ssl_session_timeout 10m;
ssl_session_tickets off;

# =====
# SECURITY HEADERS
# =====
add_header Strict-Transport-Security "max-age=31536000;
includeSubDomains" always;
add_header X-Frame-Options "SAMEORIGIN" always;
add_header X-Content-Type-Options "nosniff" always;
add_header X-XSS-Protection "1; mode=block" always;
add_header Referrer-Policy "no-referrer-when-downgrade" always;

# =====
# MAIN LOCATION - WITH RATE LIMITING
# =====
location / {
    # RATE LIMITING (BONUS 2)
    # burst=20: Allow bursts up to 20 requests
    # nodelay: Don't delay requests within burst limit
    limit_req zone=mylimit burst=20 nodelay;

    # Add rate limit information headers
    add_header X-RateLimit-Limit "10" always;

    # Proxy to load-balanced backend servers
    proxy_pass http://backend_servers;

    # Proxy Headers
    proxy_set_header Host $host;
    proxy_set_header X-Real-IP $remote_addr;
    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
    proxy_set_header X-Forwarded-Proto $scheme;
    proxy_set_header X-Forwarded-Host $host;
    proxy_set_header X-Forwarded-Port $server_port;

    # WebSocket Support
    proxy_http_version 1.1;
    proxy_set_header Upgrade $http_upgrade;
    proxy_set_header Connection "upgrade";

    # Cache Configuration
    proxy_cache my_cache;
    proxy_cache_valid 200 302 60m;
    proxy_cache_valid 404 10m;
    proxy_cache_valid 500 502 503 504 1m;
    proxy_cache_key "$scheme$request_method$host$request_uri";
}

```

```
    proxy_cache_bypass $http_cache_control $cookie_nocache
$arg_nocache;
    proxy_no_cache $httpPragma $httpAuthorization;

    # Add cache status to response headers
    add_header X-Cache-Status $upstream_cache_status always;
    add_header X-Backend-Server $upstream_addr always;

    # Timeouts
    proxy_connect_timeout 60s;
    proxy_send_timeout 60s;
    proxy_read_timeout 60s;

    # Buffers
    proxy_buffering on;
    proxy_buffer_size 4k;
    proxy_buffers 8 4k;
    proxy_busy_buffers_size 8k;
}

# =====
# HEALTH CHECK ENDPOINT
# =====
location /health {
    access_log off;
    return 200 "Nginx Load Balancer is healthy and running!\n";
    add_header Content-Type text/plain;
}

# =====
# NGINX STATUS (localhost only)
# =====
location /nginx_status {
    stub_status on;
    access_log off;
    allow 127.0.0.1;
    deny all;
}

# =====
# CUSTOM ERROR PAGES (BONUS 1)
# =====
error_page 404 /errors/404.html;
error_page 502 /errors/502.html;
error_page 503 /errors/503.html;

location /errors/ {
    internal;
    root /usr/share/nginx/html;
}
```

```

# Default error page fallback
error_page 500 504 /50x.html;
location = /50x.html {
    root /usr/share/nginx/html;
}
}

# =====
# HTTP SERVER (Port 80) - Redirect to HTTPS
# =====
server {
    listen 80;
    listen [::]:80;
    server_name _;

    # Redirect all HTTP to HTTPS
    location / {
        return 301 https://$host$request_uri;
    }

    # Health check still works over HTTP
    location /health {
        access_log off;
        return 200 "Nginx Load Balancer is healthy and running!\n";
        add_header Content-Type text/plain;
    }
}
}

```

Additional screenshots:

- All required screenshots for Parts 1–6.
- Organized in folder.

References:

- Classroom lectures or reference material.